

How Double-Fetch Situations turn into Double-Fetch Vulnerabilities: A Study of Double Fetches in the Linux Kernel

Pengfei Wang, National University of Defense Technology;

Jens Krinke, University College London;

Kai Lu and Gen Li, National University of Defense Technology;

Steve Dodier-Lazaro, University College London

26th USENIX Security Symposium,

August 16–18, 2017 • Vancouver, BC, Canada

J.T. ZHANG 2019.06, NUAA

Multicore – Concurrent Programming

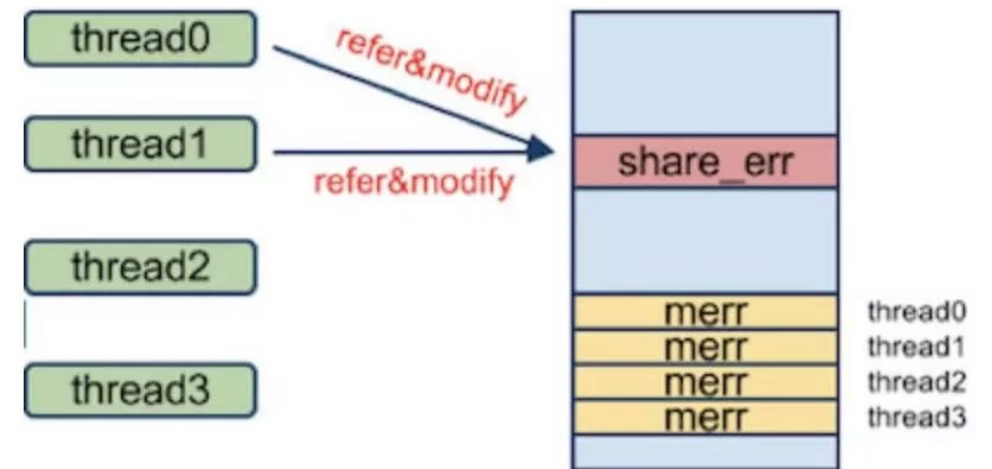
Data Race

- Shared memory
- At least one of the accesses is a write
- Atomicity-violation

```
count++;  
80485f8:  a1 2c a0 04 08      mov     0x804a02c,%eax  
80485fd:  83 c0 01             add     $0x1,%eax  
8048600:  a3 2c a0 04 08      mov     %eax,0x804a02c
```

```
1 int count = 0;
```

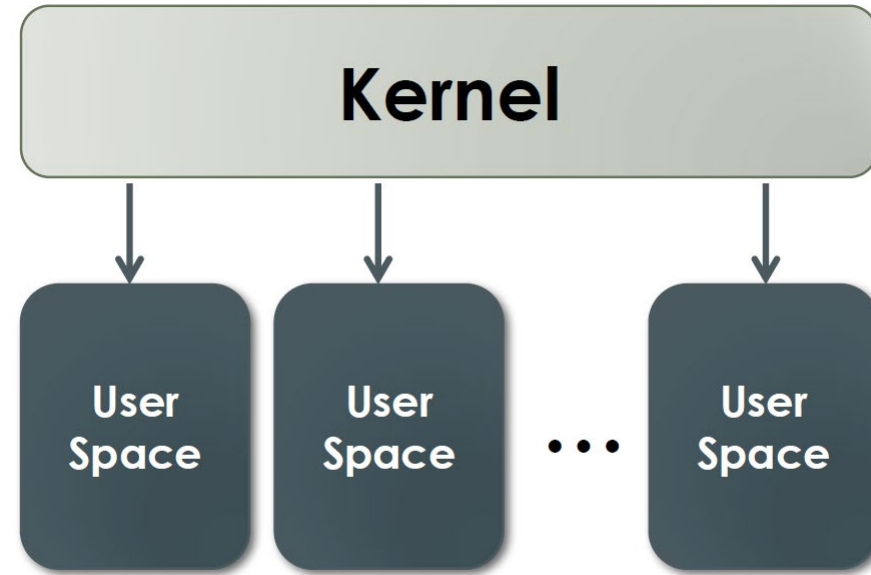
Data Race



```
1 for (unsigned int i = 0; i < 100; i++)  
2 {  
3     count++;  
4 }
```

Kernel/User Space Protection

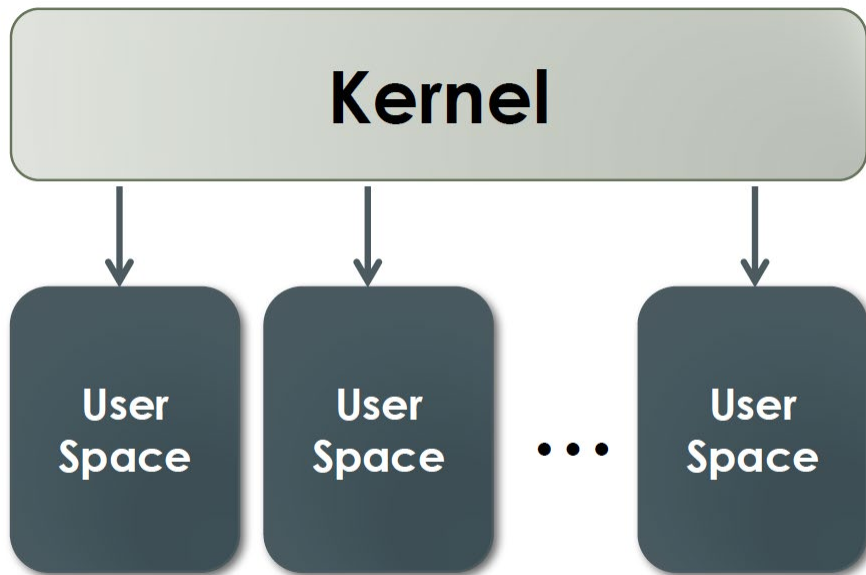
- Kernel space
- User space



- Each user process has its own virtual memory space
- User space is isolated - Virtual address → Physical address
- Only the kernel can access all user spaces

Kernel/User Space Data Exchange

- Transfer Functions



```
1 copy_from_user();  
2 copy_to_user();  
3 get_user();  
4 put_user();
```

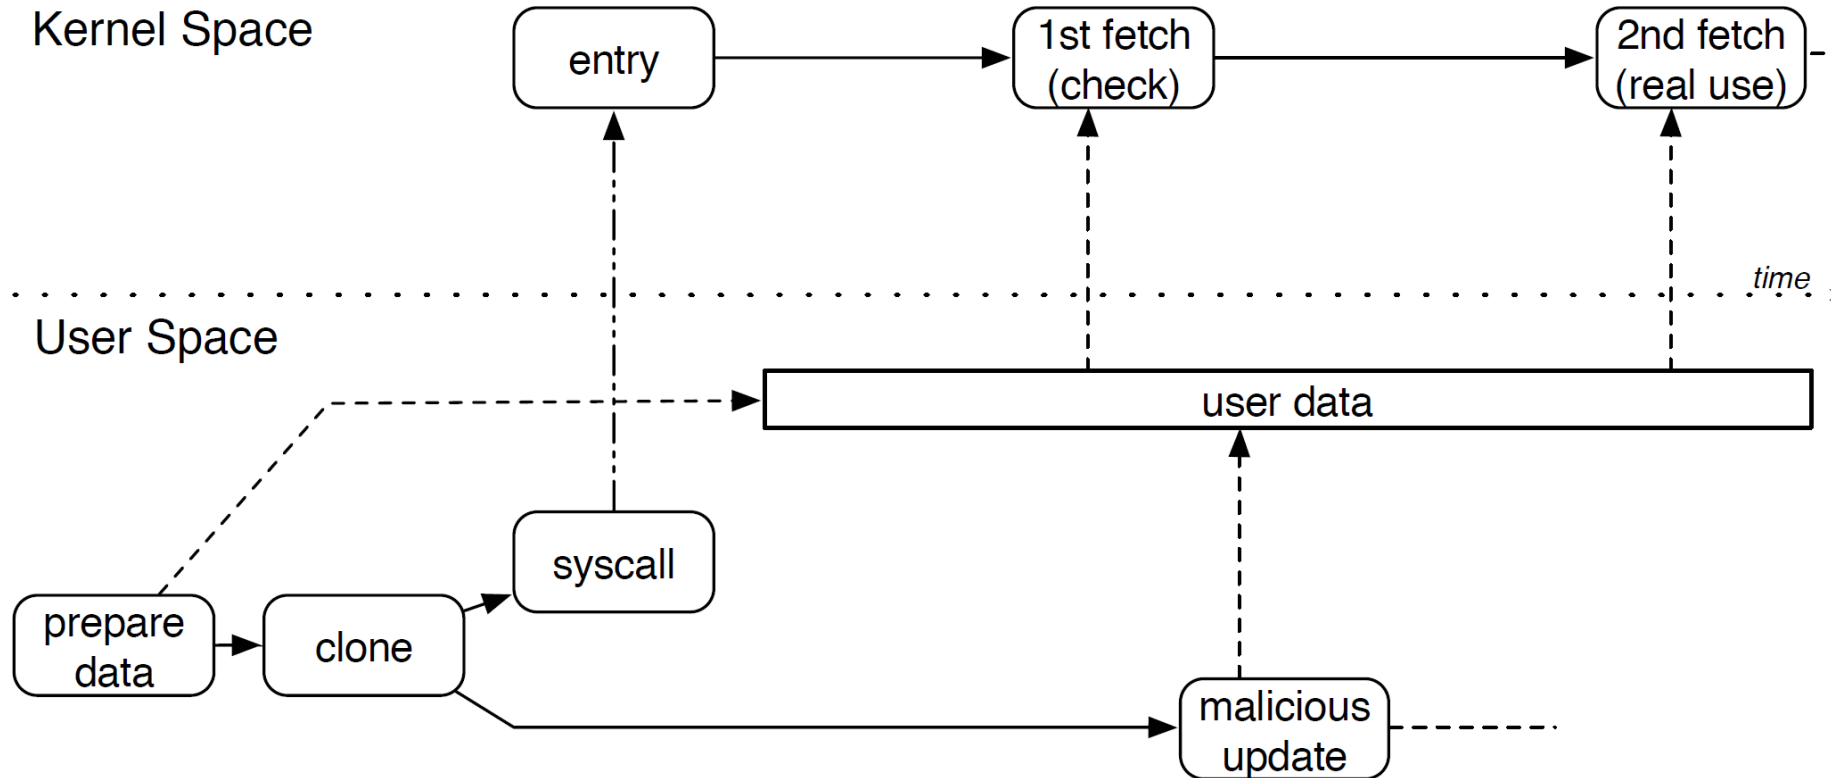
Memory Access in Drivers

- Kernel components
- Bridge between software & hardware

In Linux : Device → File

- File I/O system call
- Copy data from kernel to user space – `read();`
- Copy data from user space to kernel – `write();`

What is Double-Fetch?



- Kernel fetch data from the same memory location twice
- During two fetch, the memory is modified

Double-Fetch Grade

Benign double fetch

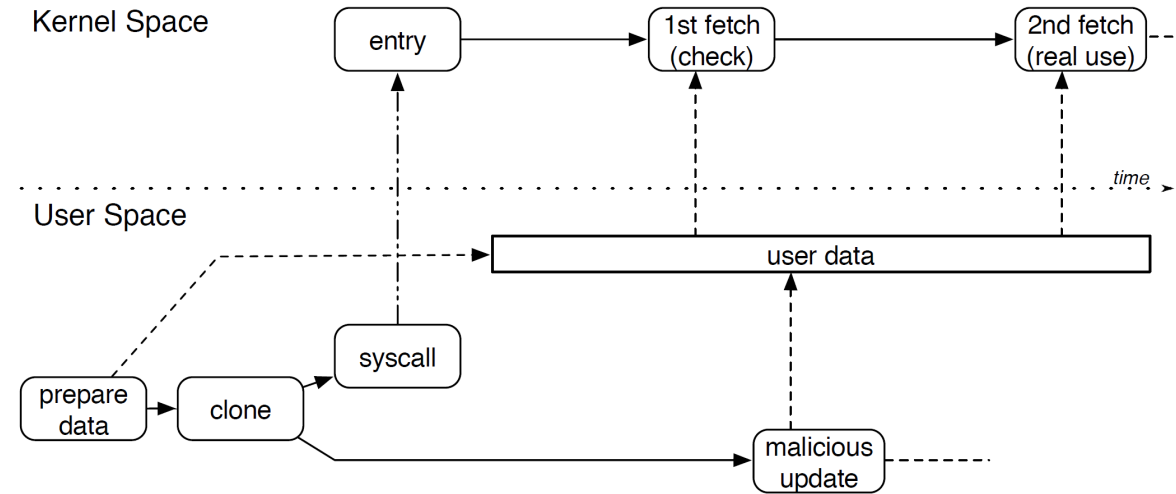
- Protection schemes
- Double-fetched value is not used

Harmful double fetch

- Could cause failures under specific conditions

Double-fetch vulnerability

- Privilege escalation
- Information leakage
- Kernel crash



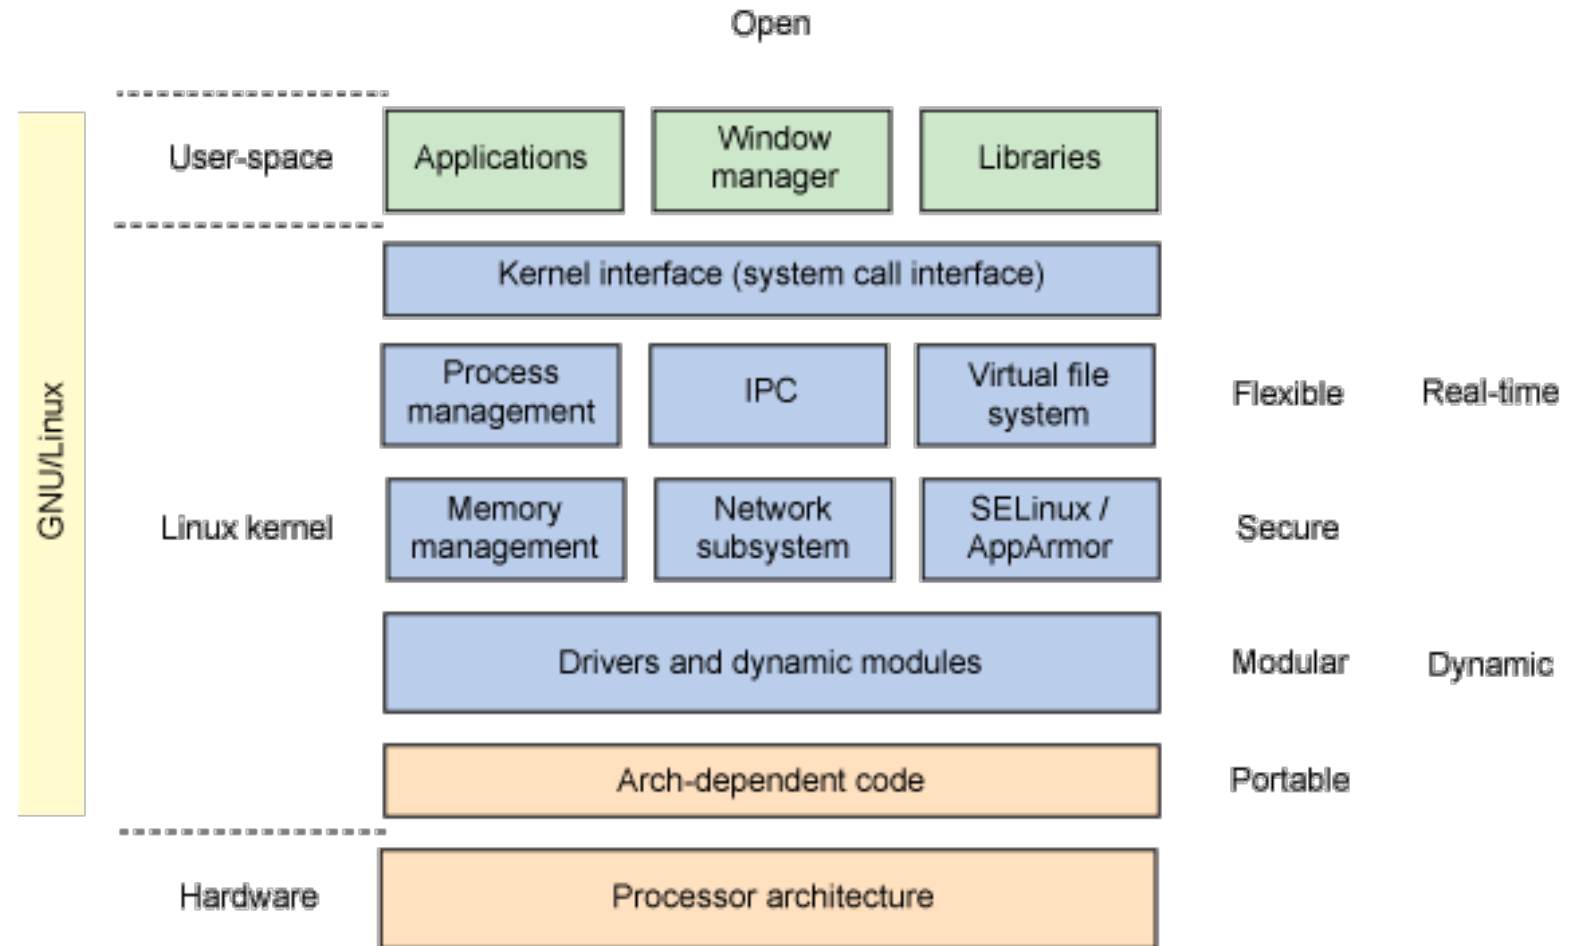
Related Works

Dynamic Analysis

- Devices?
- Architectures?
- Overhead?

Static Analysis

- Source code
- Missing context – false reports



Pattern-based static analysis

Basic Pattern Matching Analysis

- Identify multiple invocations of *transfer function*
- Analyze manually & categorize

Refined Double-Fetch Bug Detection

- Find double-fetch vulnerabilities

Basic Pattern Matching Analysis

Basic Pattern: In one function, fetching data from user space more than twice

- Matching all **get_user()** or **copy_from_user()** variants
- Target address and copy size can be different
- Source address must be the same

```
void function_name(*src)
{
    copy_from_user(dst1, src, len1)
    ...
    copy_from_user(dst2, src, len2)
}
```

Basic Pattern Matching Analysis

False Positives

- Fetching different element of the same structure
- Adding offset

Automatically removed

```
void function_name(*src)  
{  
    copy_from_user(dst1, src, len1)  
    ...  
    copy_from_user(dst2, src, len2)  
}
```


Double Fetch Categorization

- Size checking
- Type selection
- Shallow copy

Double Fetch Categorization – Size Checking

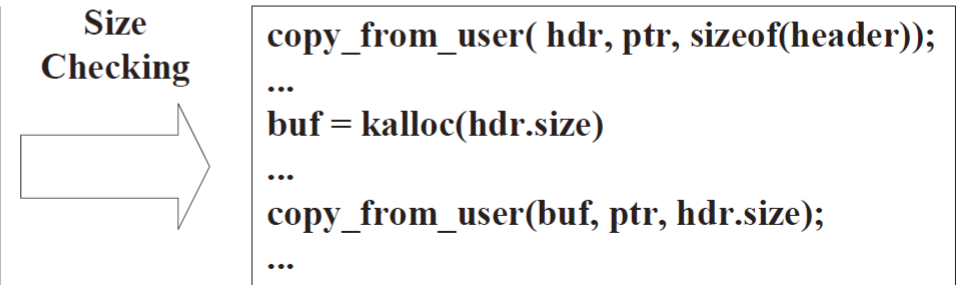
Message Data

- Header
- Body (vary in length)



First fetch

- Header only
- Get message size, check validity, allocate buffer



Second fetch

- Complete message copied to local buffer

Double Fetch Categorization – Size Checking

Adaptec RAID controller driver - commctrl.c (CVE-2016-6480)

```
60 static int ioctl_send_fib(struct aac_dev* dev,
                           void __user *arg)
61 {
62     struct hw_fib * kfib;
...
81     if (copy_from_user((void *)kfib, arg, sizeof(...))) {
82         aac_fib_free(fibptr);
83         return -EFAULT;
84     }
...
90     size = le16_to_cpu(kfib->header.Size) + sizeof(...);
91     if (size < le16_to_cpu(kfib->header.SenderSize))
92         size = le16_to_cpu(kfib->header.SenderSize);
93     if (size > dev->max_fib_size) {
...
101         kfib = pci_alloc_consistent(dev->pdev, size, &daddr);
...
114     }

115
116     if (copy_from_user(kfib, arg, size)) {
117         retval = -EFAULT;
118         goto cleanup;
119     }
120
121     if (kfib->header.Command == cpu_to_le16(...)) {
...
128     } else {
129         retval =
            aac_fib_send(le16_to_cpu(kfib->header.Command), ...
130                        le16_to_cpu(kfib->header.Size) , FsaNormal,
131                        1, 1, NULL, NULL);
...
139     }
...
160 }
```

Double Fetch Categorization – Size Checking

Results

- 30 occurrences
- 22 in drivers
- 4 of 22 in drivers occurrences cause double-fetch bugs

Double Fetch Categorization – Type Selection

Message Data

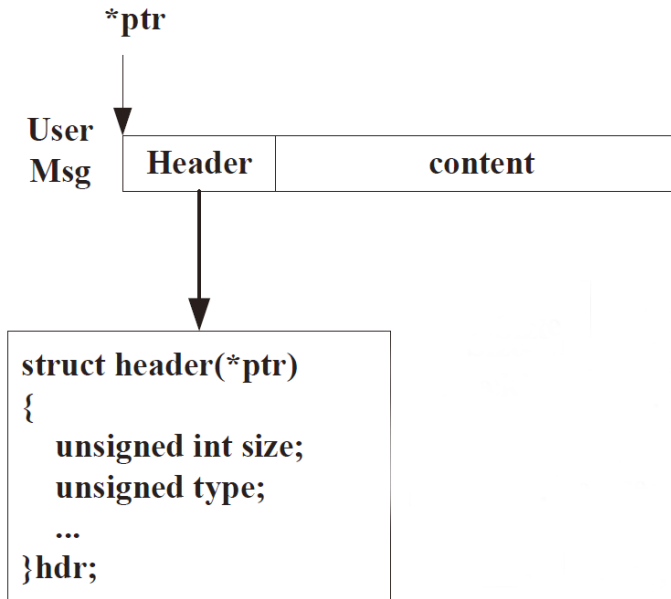
- Header
- Body

First fetch

- Header only
- Recognize message type

Second fetch

- Complete message copied to local buffer



Type
Selection

```
copy_from_user( hdr, ptr, sizeof(header));

switch(hdr.type){
    case 1:
        copy_from_user()
        ...
    case 2:
        copy_from_user()
        ...
    default:
        ...
}
```


Double Fetch Categorization – Type Selection

Network driver - xgb3_main.c

No vulnerability

```
2129 static int cxgb_extension_ioctl(struct net_device *dev,  
                                   void __user *useraddr)
```

```
2130 {
```

```
...
```

```
2133     u32 cmd;
```

```
...
```

```
2136     if (copy_from_user(&cmd, useraddr, sizeof(cmd)))
```

```
2137         return -EFAULT;
```

```
2138
```

```
2139     switch (cmd) {
```

```
2140     case CHELSIO_SET_QSET_PARAMS:{
```

```
...
```

```
2143     struct ch_qset_params t;
```

```
...
```

```
2149     if (copy_from_user(&t, useraddr, sizeof(t)))
```

```
2150         return -EFAULT;
```

```
2151     if (t.qset_idx >= SGE_QSETS)
```

```
2152         return -EINVAL;
```

```
...
```

```
2238     break;
```

```
2239 }
```

```
2284     case CHELSIO_SET_QSET_NUM:{
```

```
2285         struct ch_reg edata;
```

```
...
```

```
2292     if (copy_from_user(&edata, useraddr, sizeof(edata)))
```

```
2293         return -EFAULT;
```

```
2294     if (edata.val < 1 ||
```

```
2295         (edata.val > 1 && !(...)))
```

```
2296         return -EINVAL;
```

```
...
```

```
2313     break;
```

```
2314 }
```

Double Fetch Categorization – Type Selection

Results

- 11 occurrences
- 9 in drivers
- None of 11 occurrences cause double-fetch bugs

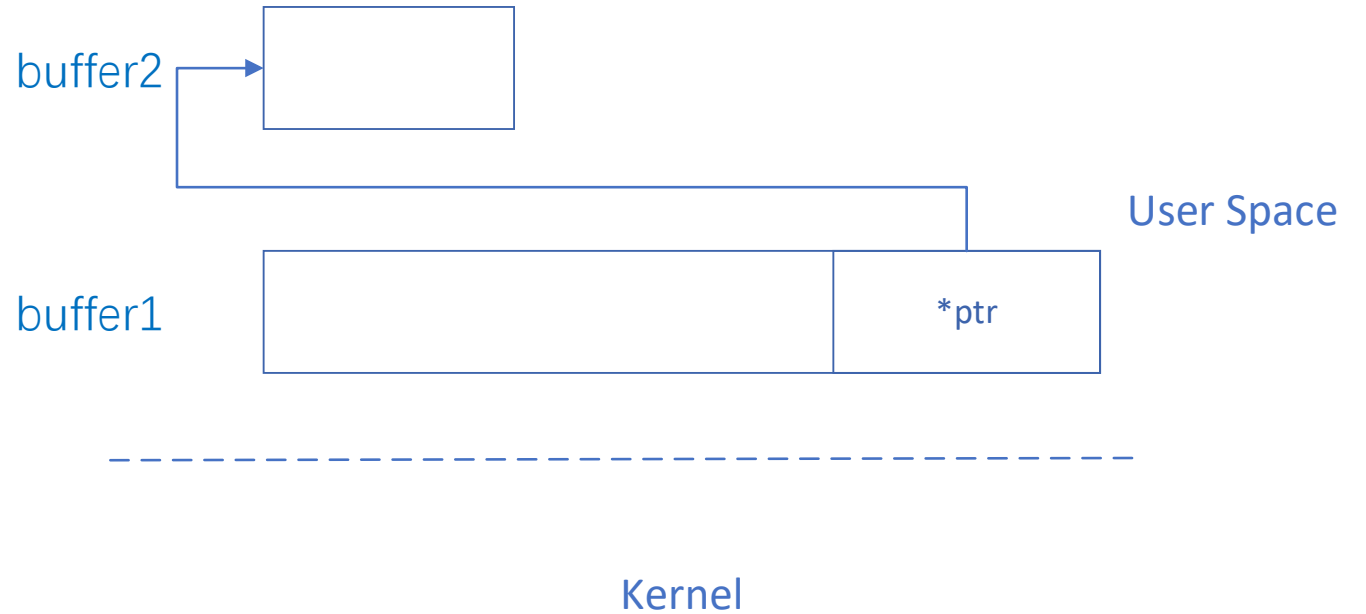
Double Fetch Categorization – Shallow Copy

Copy buffer1 → Shallow copy

Copy buffer2 → Deep copy

Results

- 31 occurrences
- 19 in drivers



Double Fetch Categorization – Shallow Copy

IBM S/390 SCLP console driver - sclp_ctl.c (CVE-2016-6130)

```
55 static int sclp_ctl_ioctl_sccb(void __user *user_area)
56 {
57     struct sclp_ctl_sccb ctl_sccb;
58     struct sccb_header *sccb;
59     int rc;
60
61     if (copy_from_user(&ctl_sccb, user_area,
62                        sizeof(ctl_sccb)))
63         return -EFAULT;
64     ...
65     sccb = (void *) get_zeroed_page(GFP_KERNEL | GFP_DMA);
66     if (!sccb)
67         return -ENOMEM;
68     if (copy_from_user(sccb, u64_to_uptr(ctl_sccb.sccb),
69                        sizeof(*sccb))) {
70         rc = -EFAULT;
71         goto out_free;
72     }
73     if (sccb->length > PAGE_SIZE || sccb->length < 8)
74         return -EINVAL;
75     if (copy_from_user(sccb, u64_to_uptr(ctl_sccb.sccb),
76                        sccb->length)) {
77         rc = -EFAULT;
78         goto out_free;
79     }
80     ...
81     if (copy_to_user(u64_to_uptr(ctl_sccb.sccb), sccb,
82                     sccb->length))
83         rc = -EFAULT;
84     ...
85 }
86 }
```

Refined Double-Fetch Bug Detection

- Improved analysis
- Specifically identify double-fetch bugs

Rule 0 (basic rule) – two reads fetch data from the same location

- More rules...
- Coccinelle applied these rules one by one

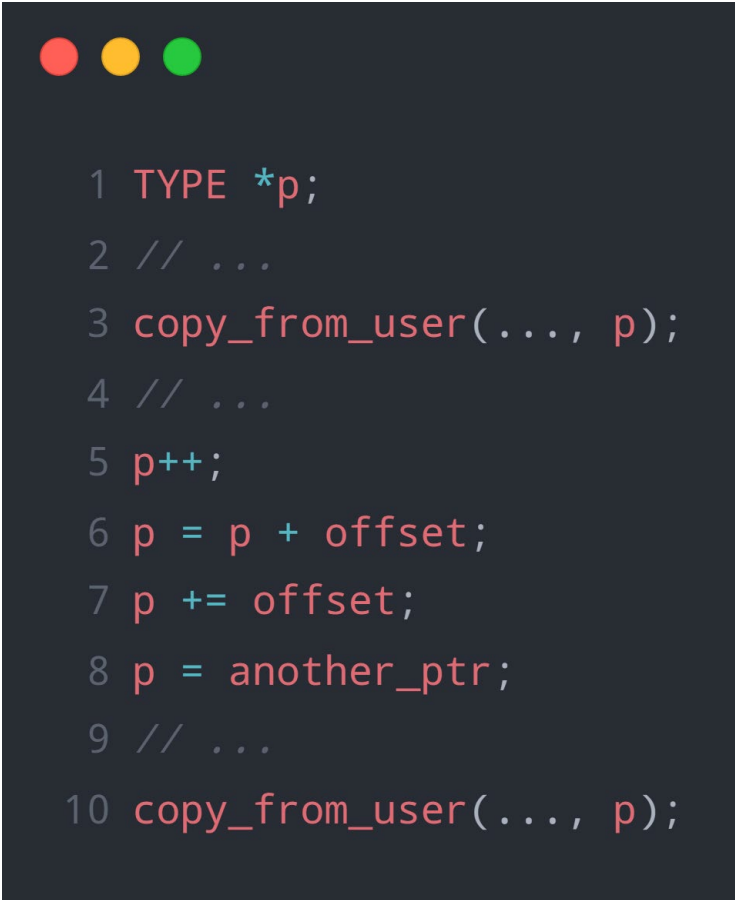
Refined Double-Fetch Bug Detection

Rule 1 – No pointer change

Keeping the user pointer unchanged between two fetches

- Self-increment (++/--)
- Adding an offset
- Assignment of another value

↓ False positives



```
1 TYPE *p;  
2 // ...  
3 copy_from_user(..., p);  
4 // ...  
5 p++;  
6 p = p + offset;  
7 p += offset;  
8 p = another_ptr;  
9 // ...  
10 copy_from_user(..., p);
```

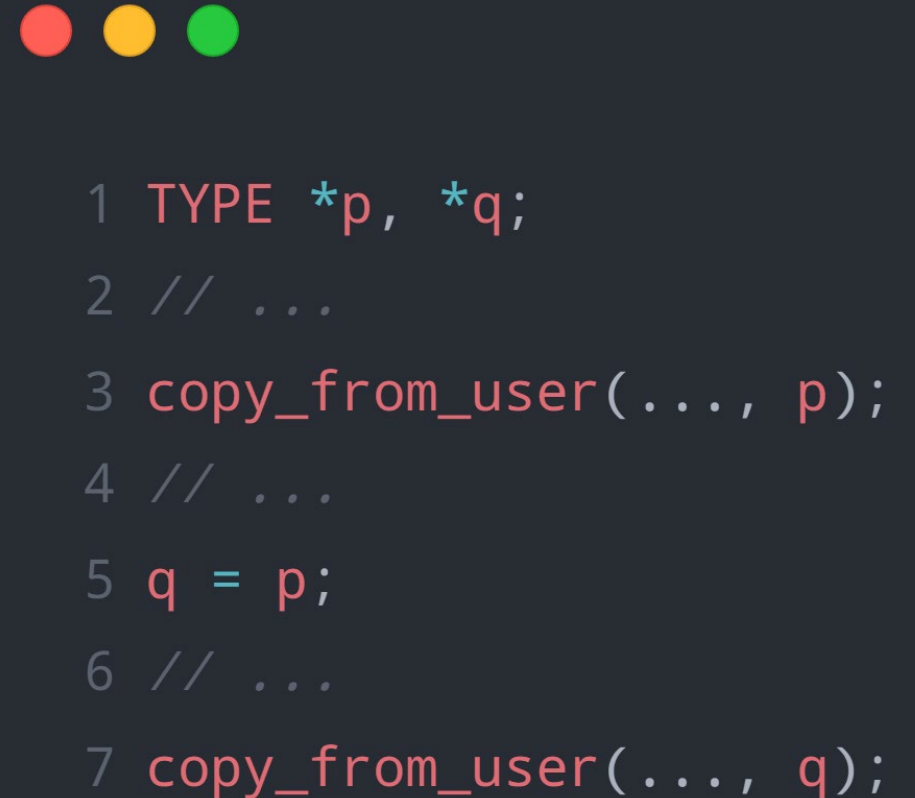
Refined Double-Fetch Bug Detection

Rule 2 – Pointer Aliasing

Source pointer is assigned to another pointer

- For convenience

↓ False negatives



```
1 TYPE *p, *q;  
2 // ...  
3 copy_from_user(..., p);  
4 // ...  
5 q = p;  
6 // ...  
7 copy_from_user(..., q);
```


Refined Double-Fetch Bug Detection

Rule 3 – Explicit type conversion

Causing the same memory region to be manipulated by two types of pointers

- First fetch – (header pointer)
- Second fetch – message pointer

↓ False negatives



```
1 TYPE *p;  
2 // ...  
3 copy_from_user(..., p);  
4 // ...  
5 copy_from_user(..., (AnotherType*)p);
```



Refined Double-Fetch Bug Detection

Rule 4 – Combination of element fetch and pointer fetch

Pointer is both used to fetch the **whole data structure** as well as only **a part**

- Not using the same pointer
- But cover the same value

↓ False negatives



```
1 TYPE *p;  
2 // ...  
3 copy_from_user(..., p->len);  
4 // ...  
5 copy_from_user(..., p);
```

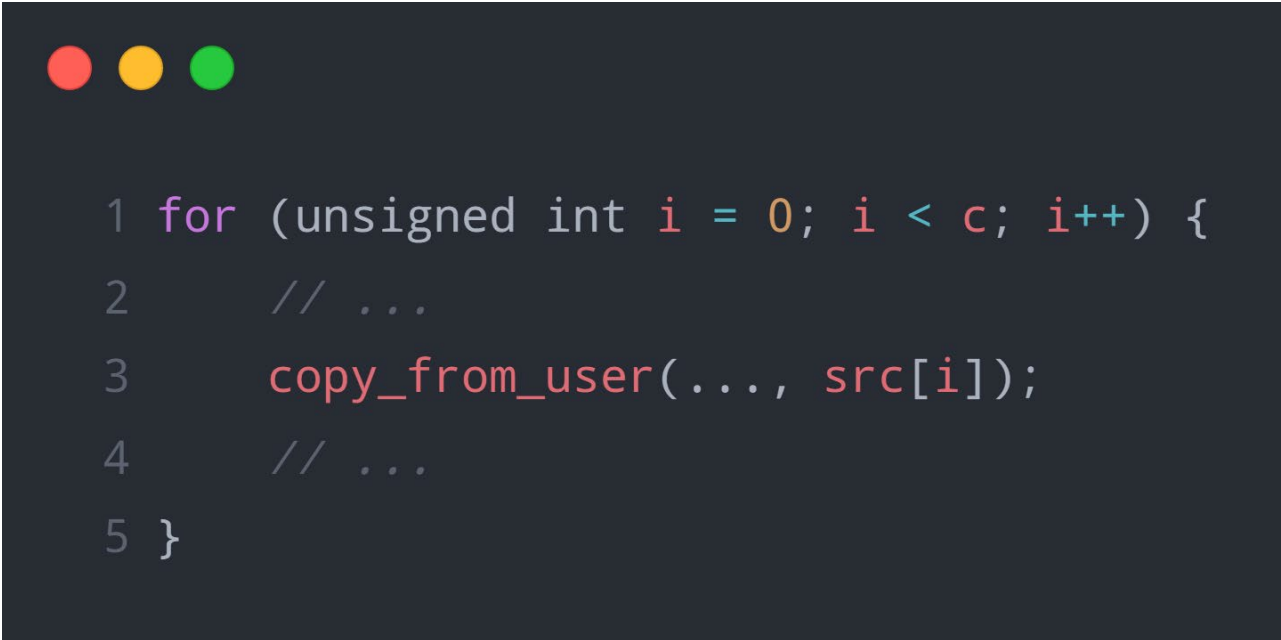
Refined Double-Fetch Bug Detection

Rule 5 – Loop involvement

One transfer function call in a loop will be reported as two calls

- Should be removed
- Cross-iteration transfer function call

↓ False positives




```
1 for (unsigned int i = 0; i < c; i++) {  
2     // ...  
3     copy_from_user(..., src[i]);  
4     // ...  
5 }
```

Tool - Coccinelle

An open-source utility for matching and transforming the source code of programs written in the C programming language using a "semantic patch" syntax based on the patch syntax.

Semantic Patch Language (SmPL)



Evaluation

- Linux 4.5
- Android 6.0.1 (based on Linux 3.18)
- FreeBSD (master branch)

- Drivers are the hard-hit area
- Benign case could turn into a vulnerability with code update (CVE-2016-5728)

Double-Fetch BUG Prevention

- Don't copy the header twice (only copy body in the second fetch)
- Use the same value
 - Only use the data from one of the fetches
- Overwrite data
 - Overwrite the header with the first fetch (widely adopted in FreeBSD)
- Compare data
 - Compare before using, abort the operation if the data is not the same
- Synchronize fetches
 - Guarantee the atomicity
 - Performance penalties

Limitations

- Preprocessed code
- Compiled code

Double-fetch can occur in:

- Macros
- Compiler optimization (in compiled binary, not in source code)
- Not labeled as *volatile*
 - Turn the memory access from single to multiple at the binary level

Thanks.

Q&A ?