


Precise and Scalable Detection of Double-Fetch Bugs in OS Kernels

Meng Xu, Chenxiong Qian, Kangjie Lu⁺, Michael Backes*, Taesoo Kim

*Georgia Tech | University of Minnesota⁺ | CISPA, Germany**

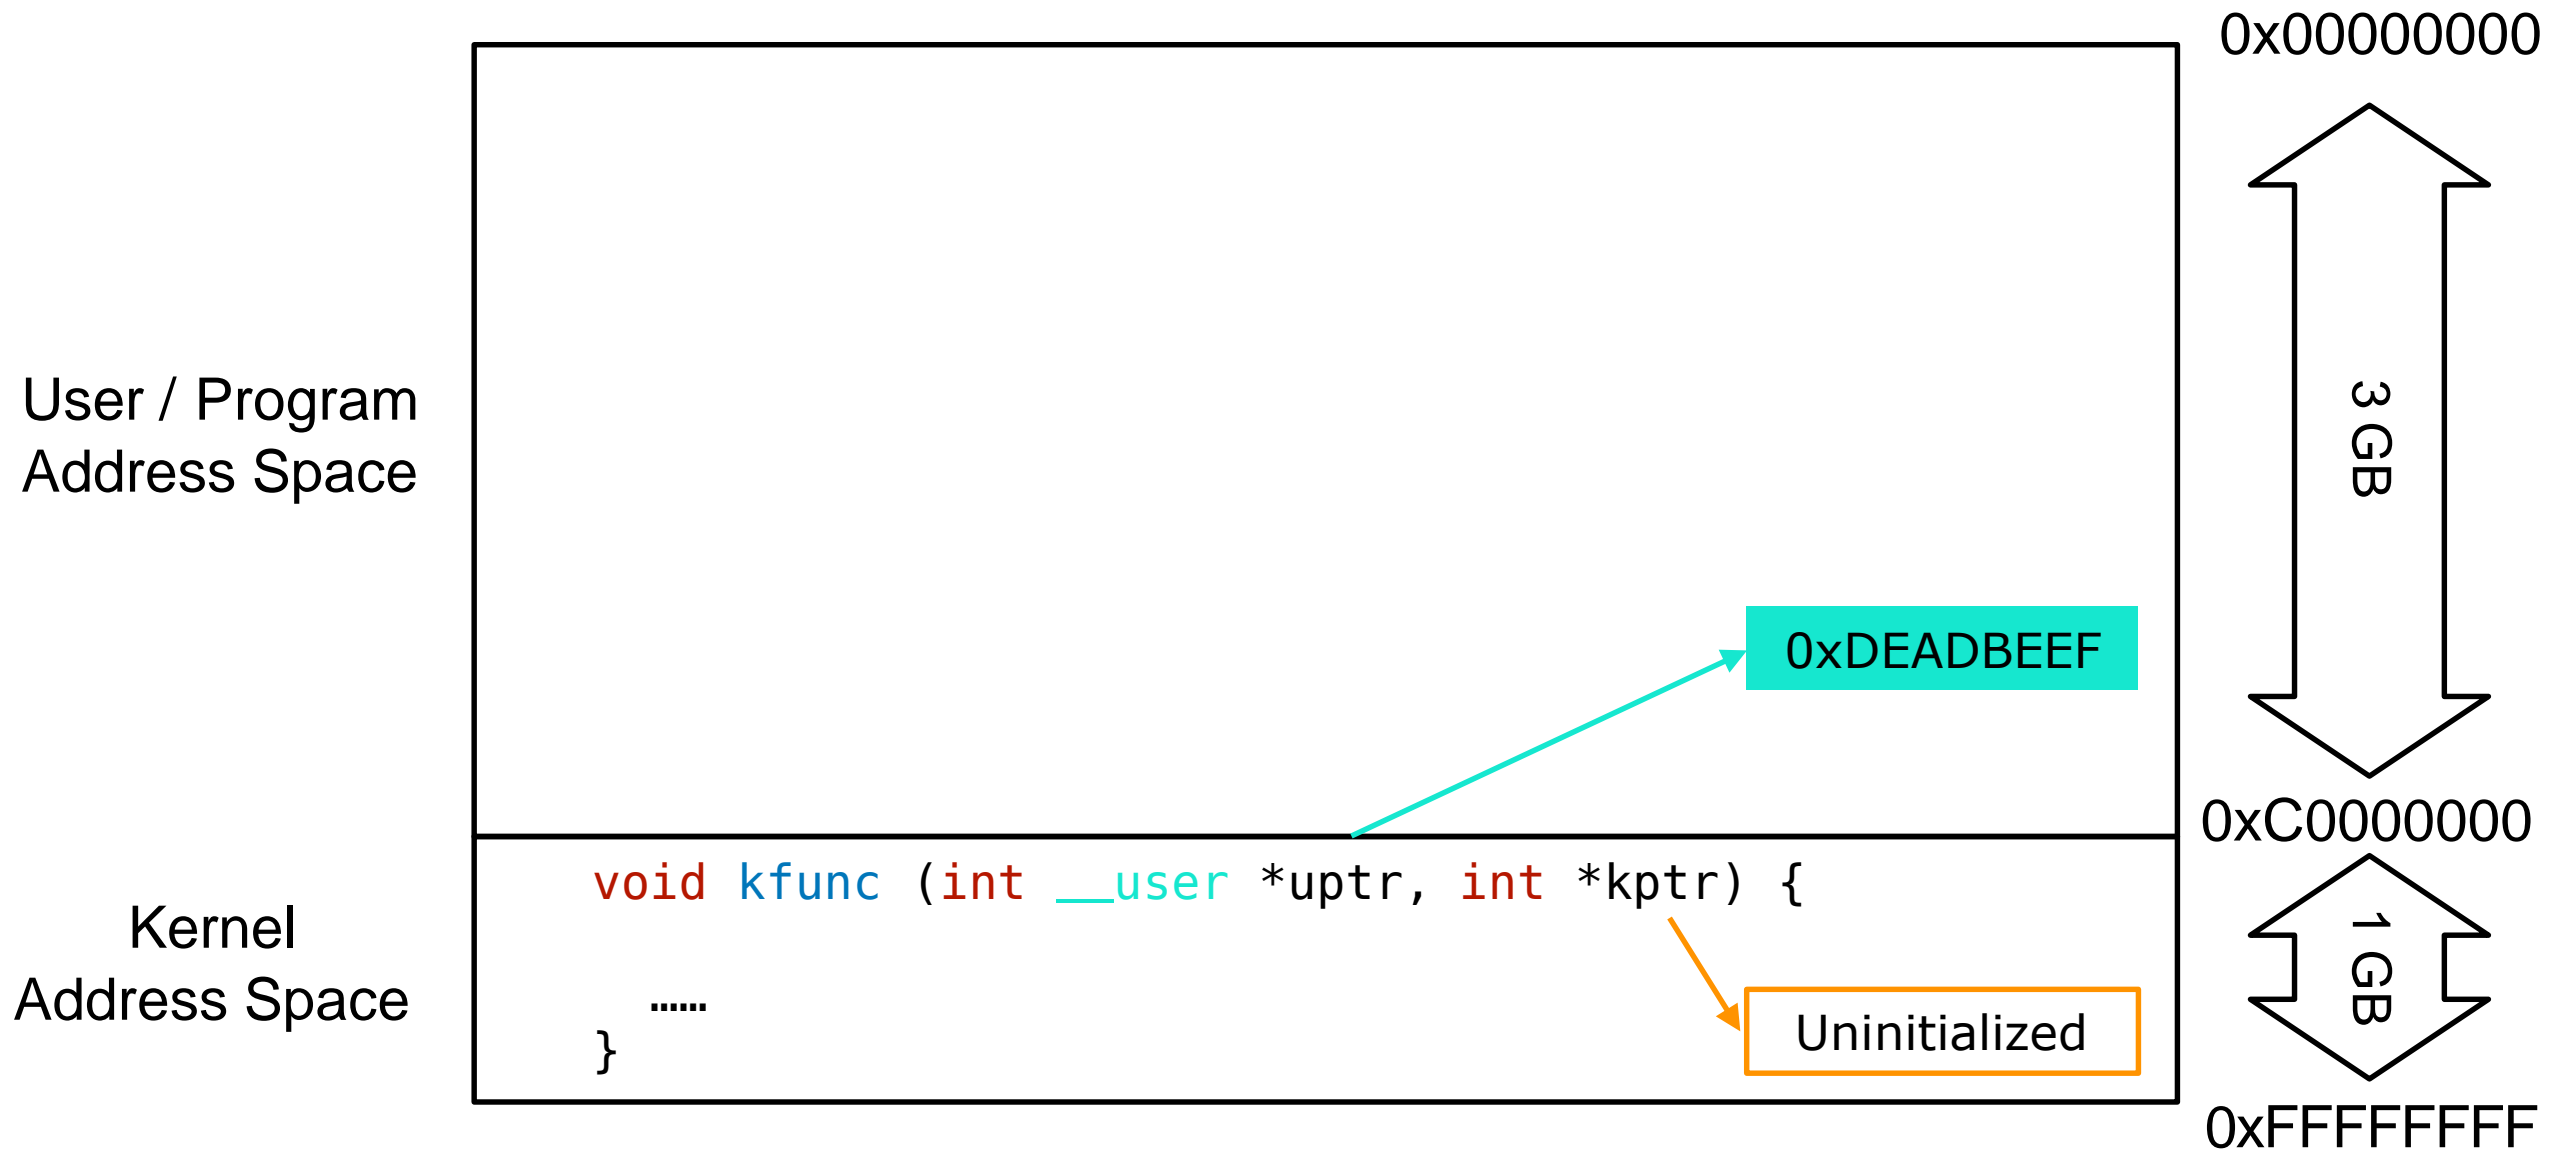
What is Double-Fetch?

Address Space Separation



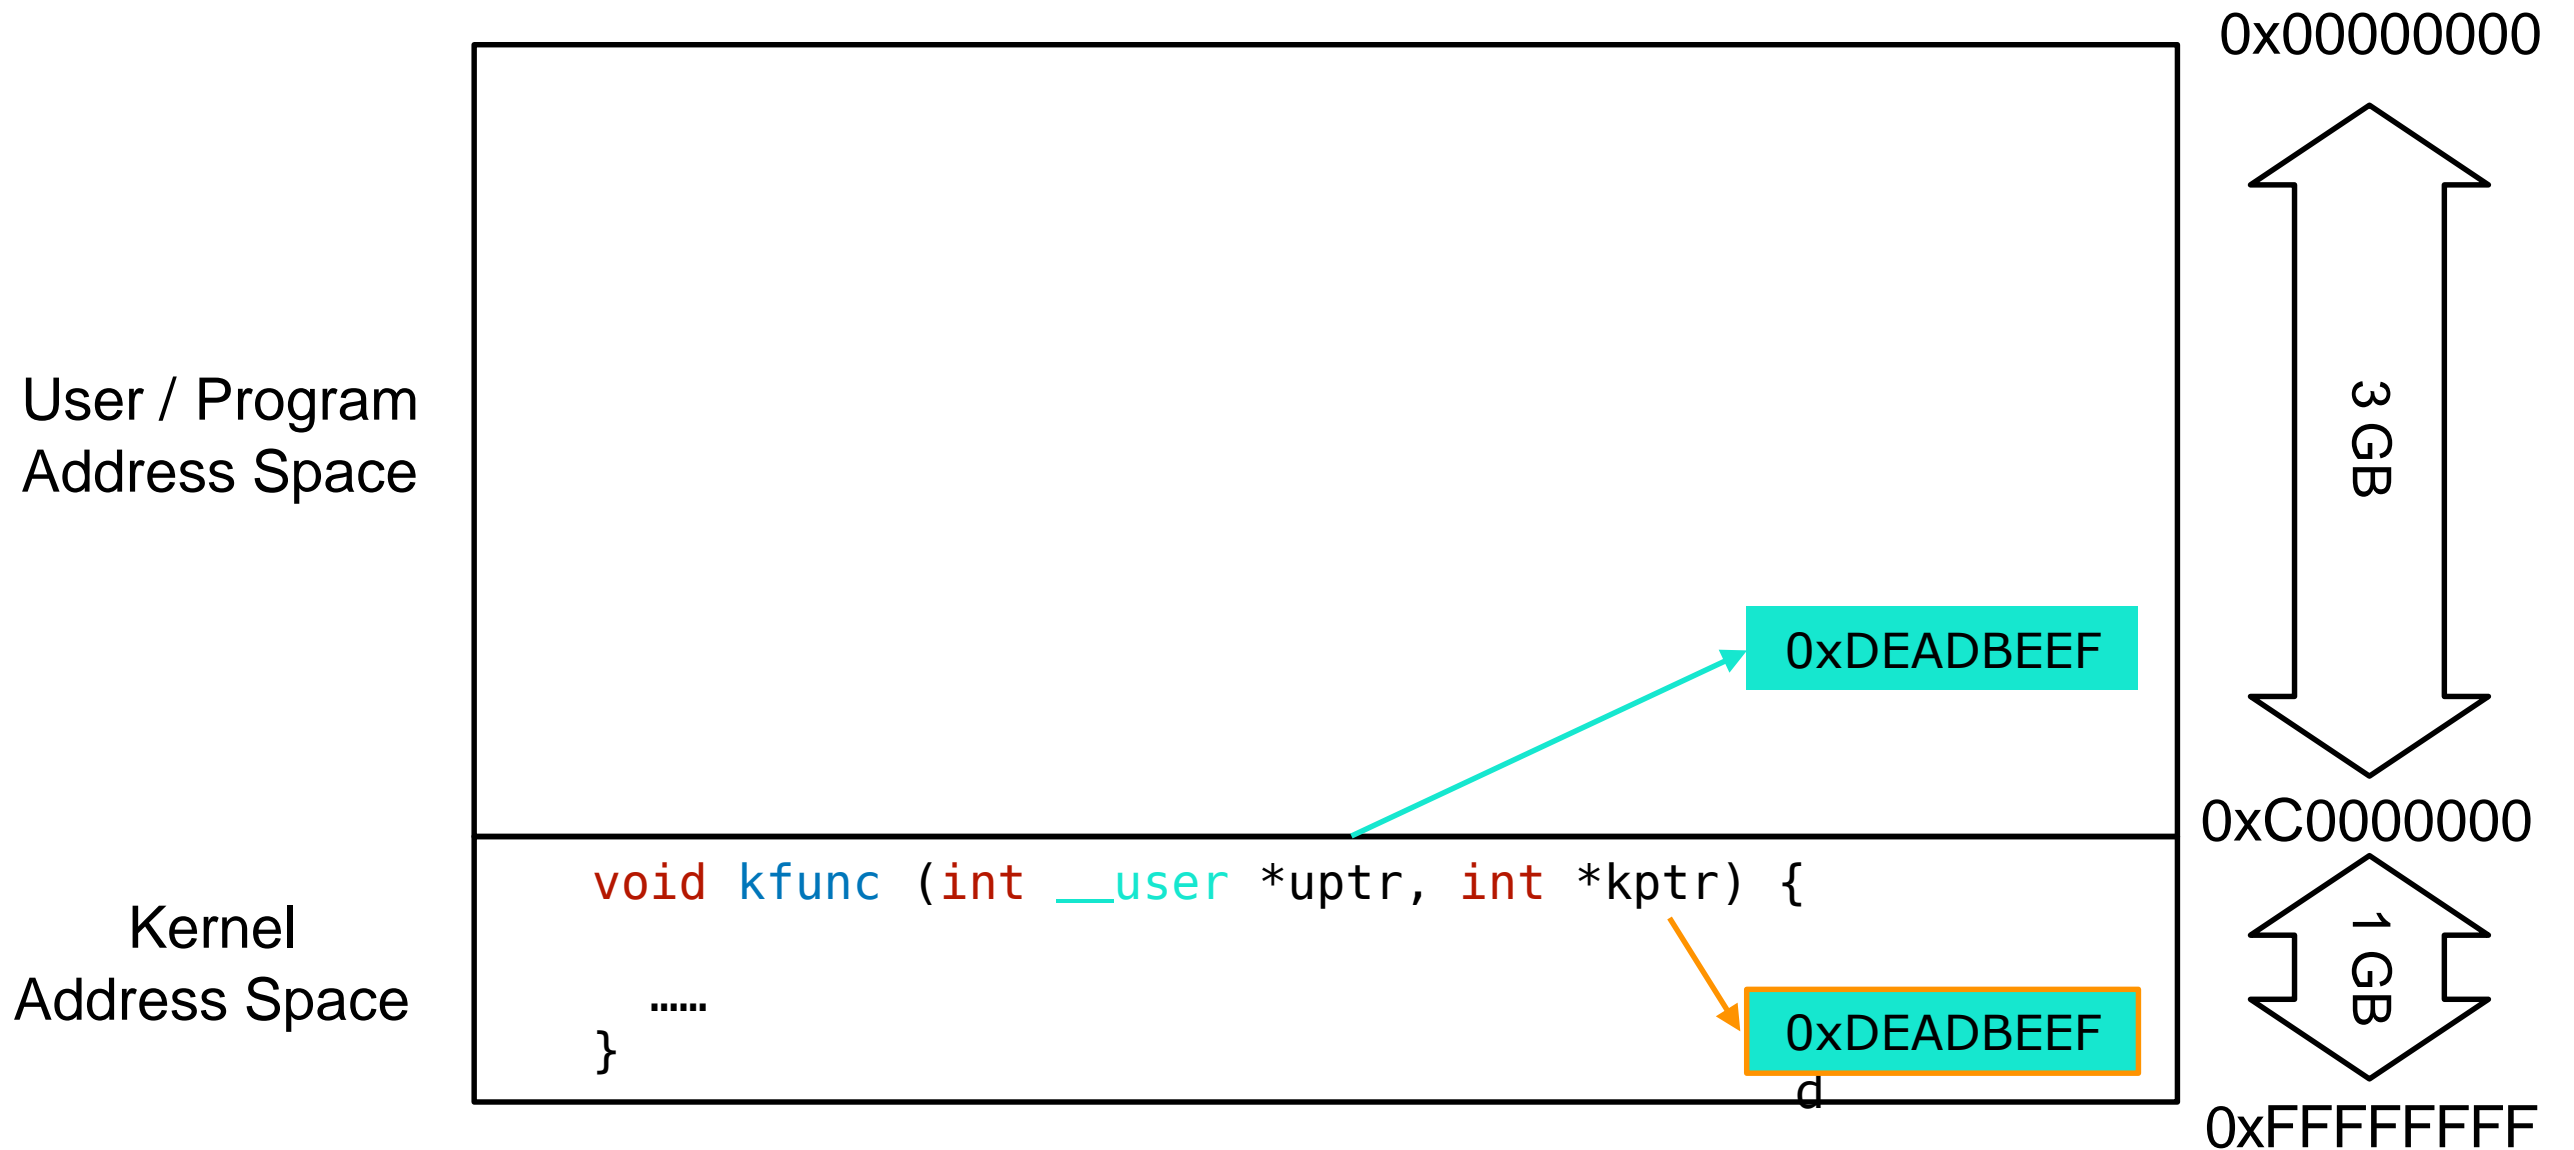
A Typical Address Space Separation Scheme with a 32-bit Virtual Address Space

No Dereference on Userspace Pointers



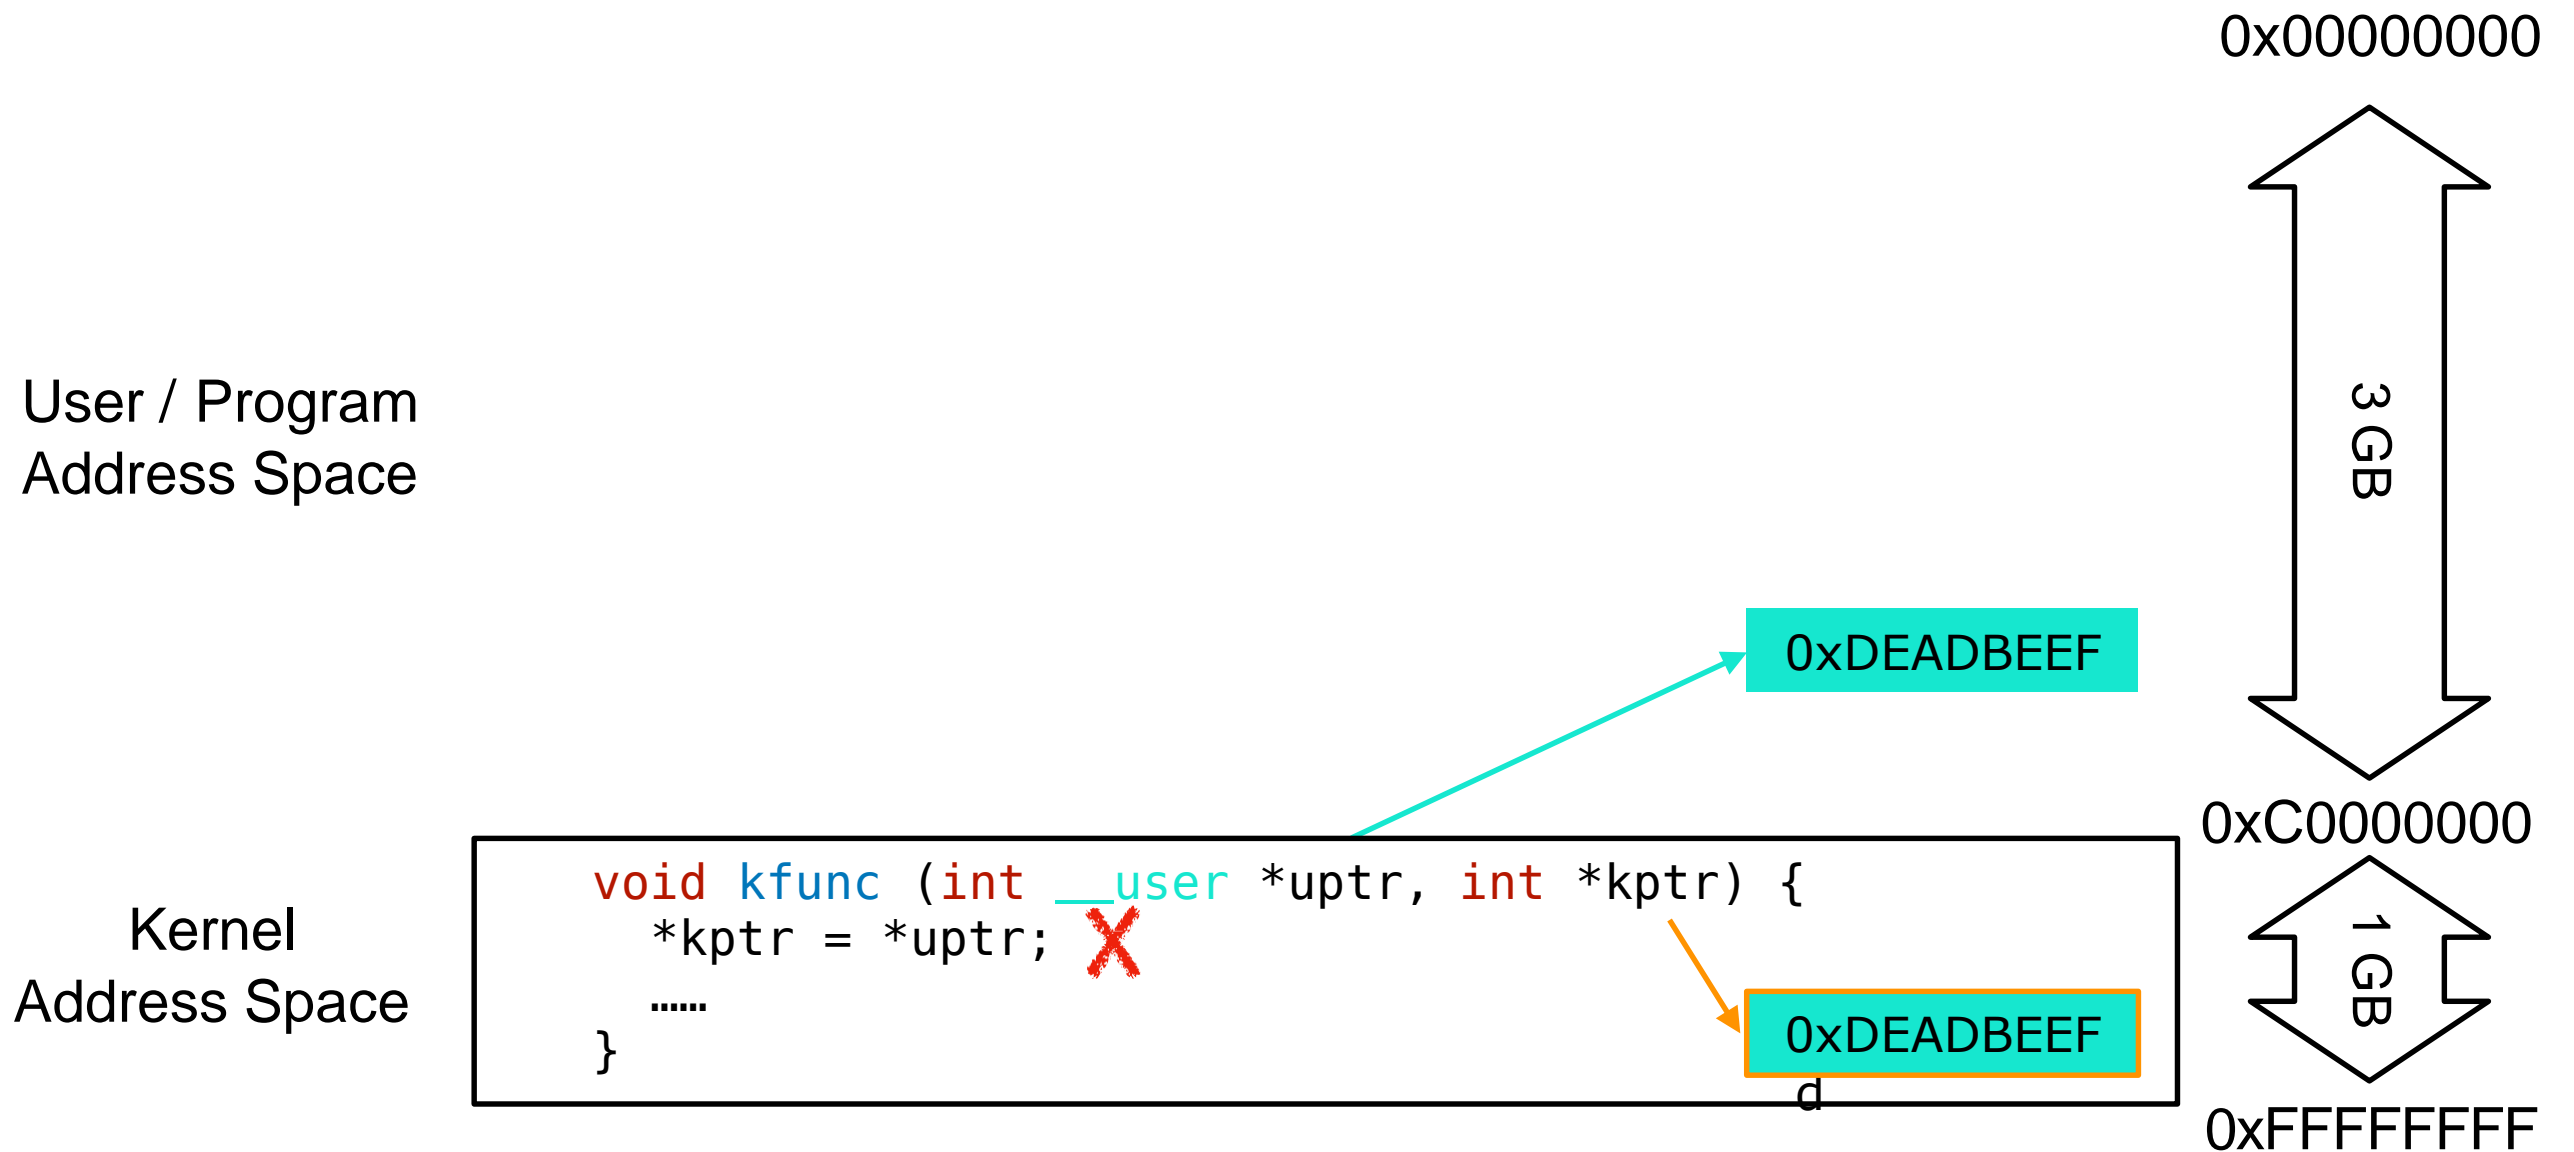
A Typical Address Space Separation Scheme with a 32-bit Virtual Address Space

No Dereference on Userspace Pointers



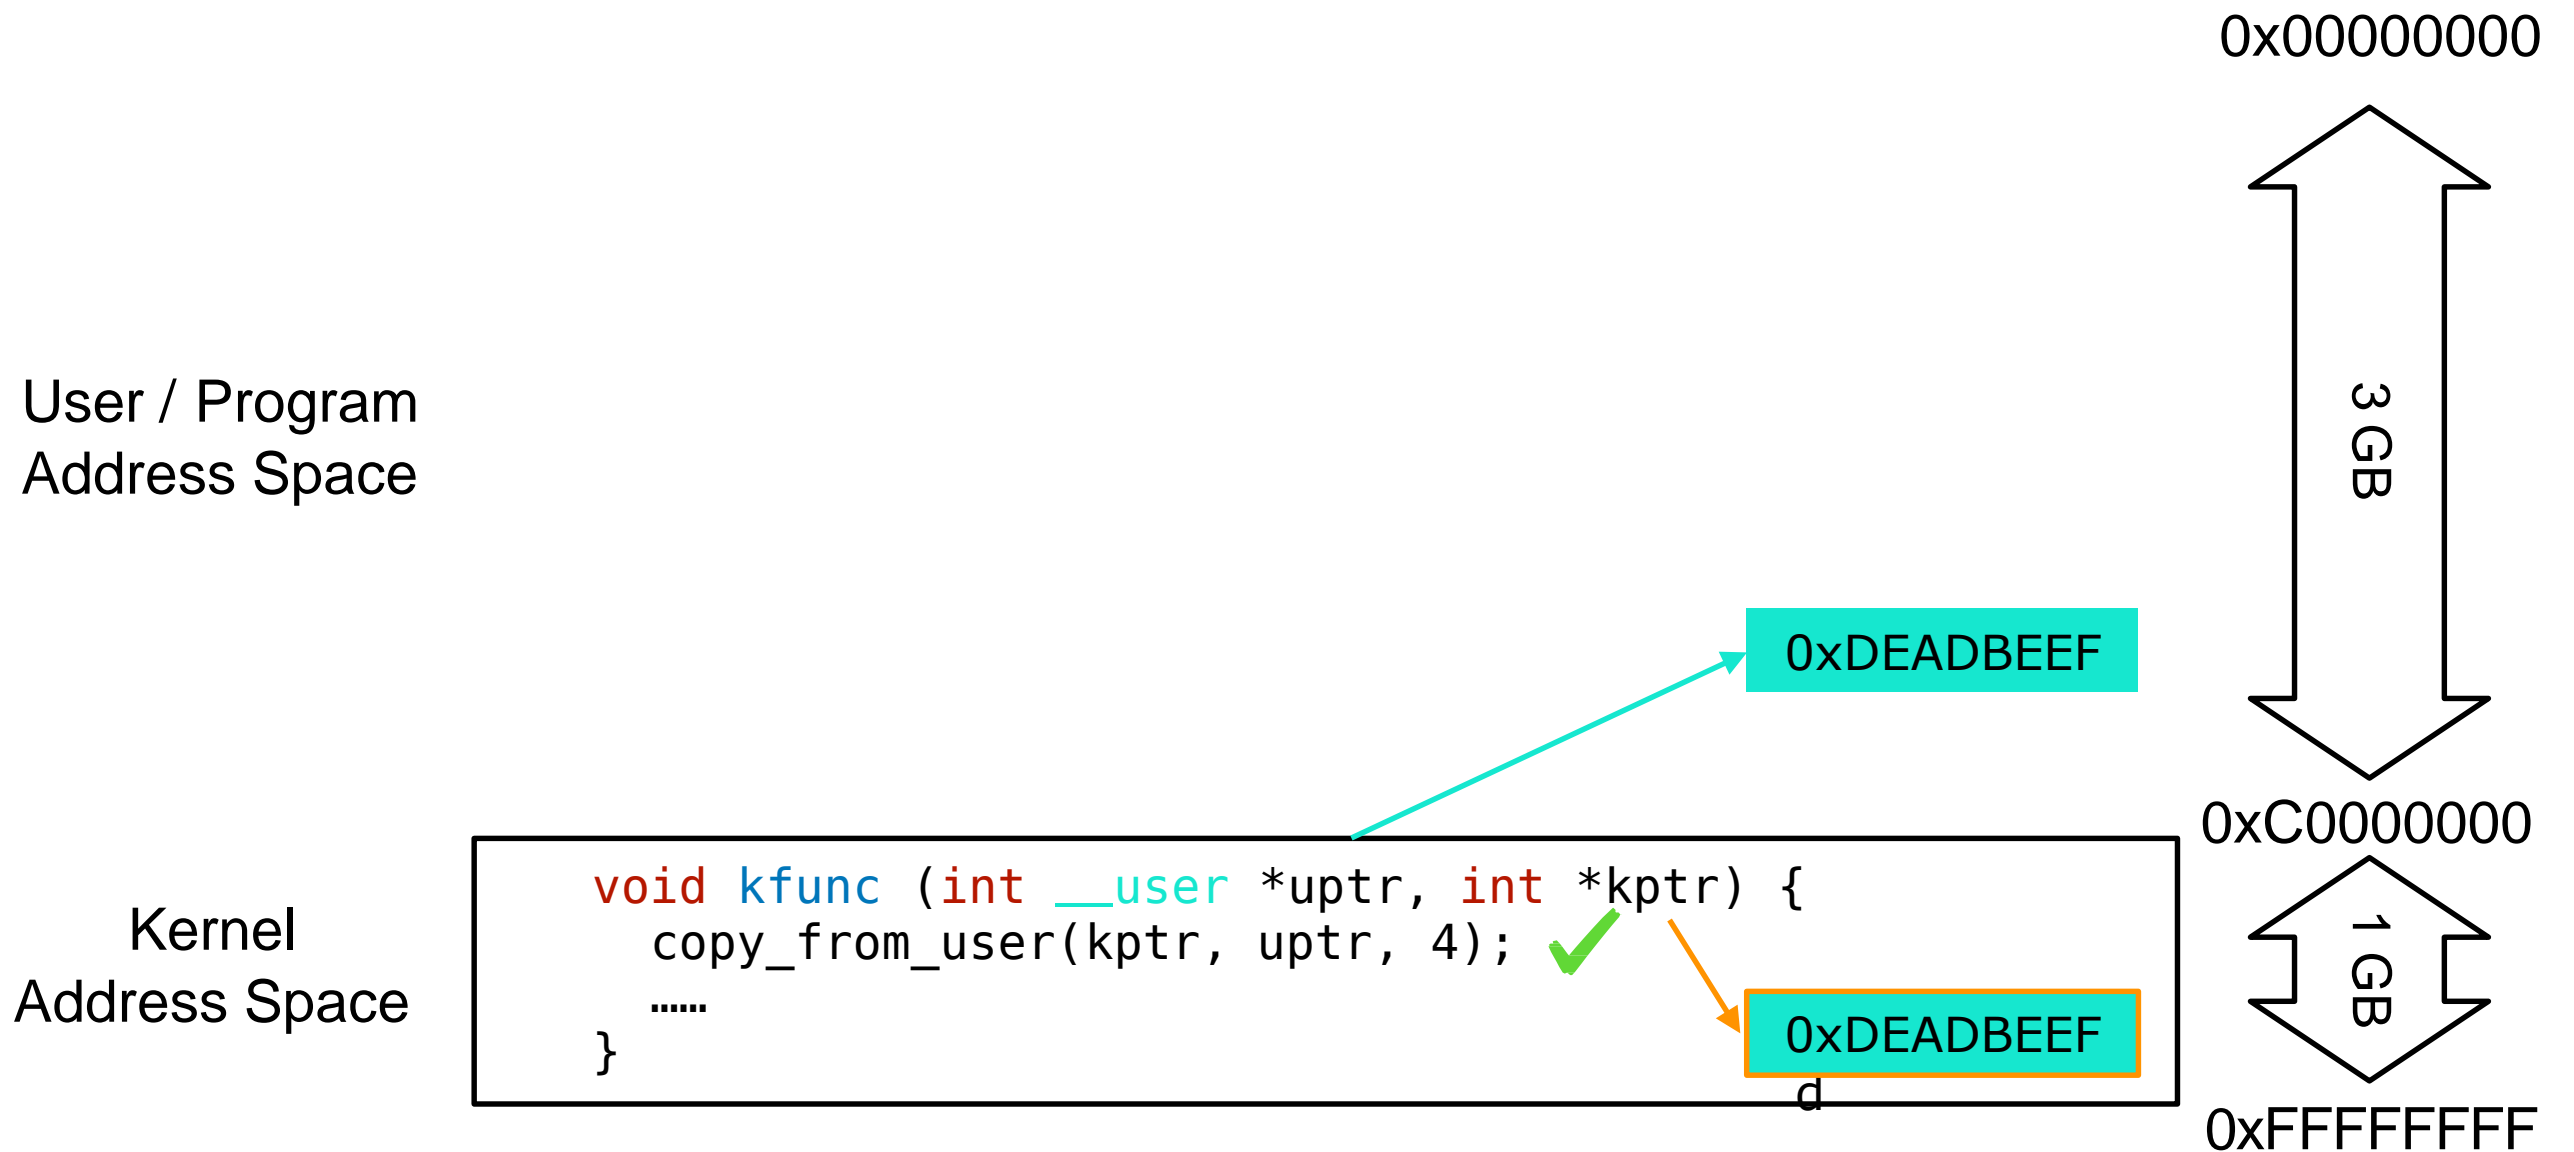
A Typical Address Space Separation Scheme with a 32-bit Virtual Address Space

No Dereference on Userspace Pointers



A Typical Address Space Separation Scheme with a 32-bit Virtual Address Space

No Dereference on Userspace Pointers



A Typical Address Space Separation Scheme with a 32-bit Virtual Address Space

Why Double-Fetch?

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
```


?? bytes

```
24 memcpy(buf, attr, attr->size);
```

Adapted from `perf_copy_attr` in file `kernel/events/core.c`

Why Double-Fetch?

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
```




```
24 memcpy(buf, attr, attr->size);
```

Adapted from `perf_copy_attr` in file `kernel/events/core.c`

Why Double-Fetch?

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9     return -EFAULT;
```



```
24 memcpy(out, attr, attr->size);
```

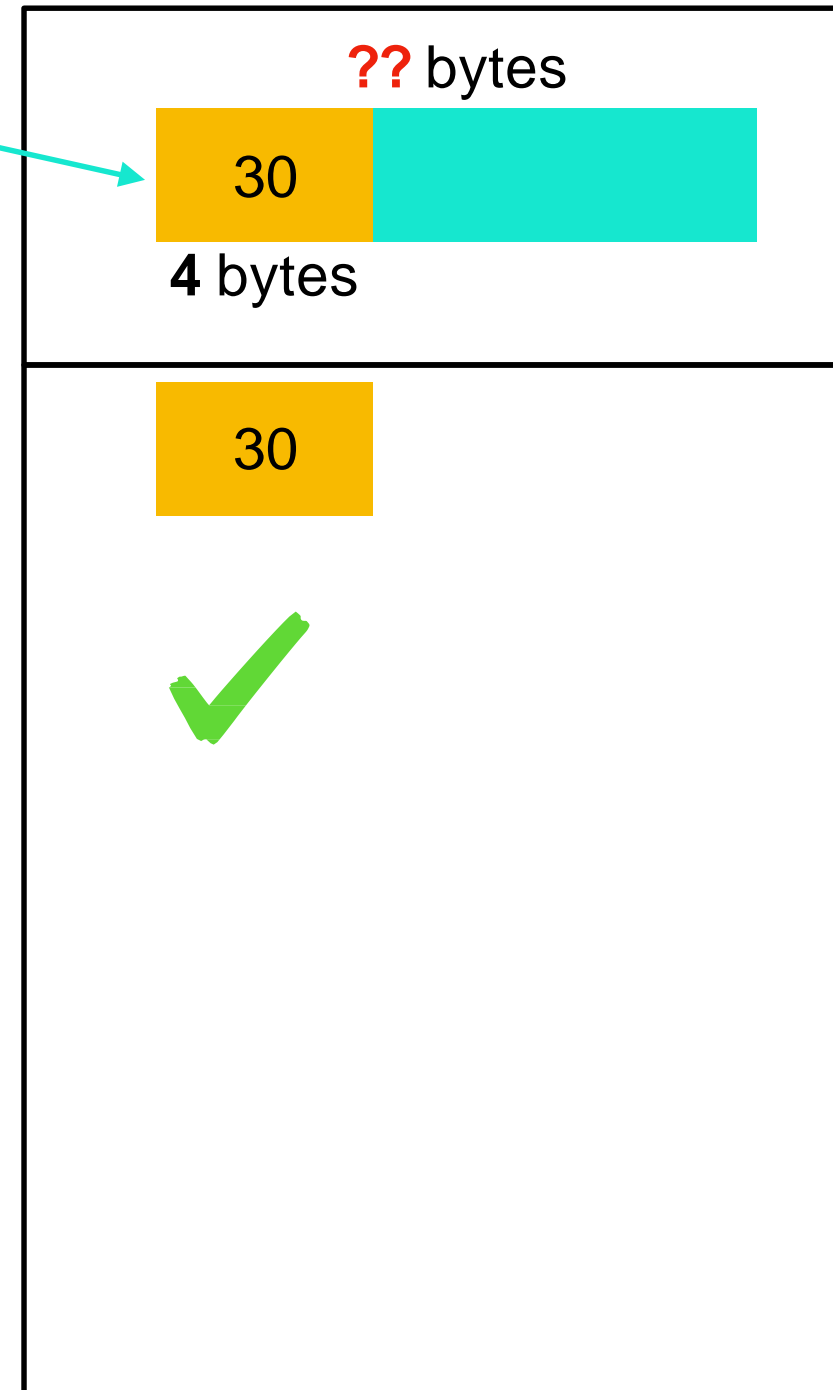
Adapted from `perf_copy_attr` in file `kernel/events/core.c`

Why Double-Fetch?

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9     return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size < PERF_ATTR_SIZE_VER0)
14    return -EINVAL;
```

```
24 memcpy(out, attr, attr->size);
```

Adapted from `perf_copy_attr` in file `kernel/events/core.c`

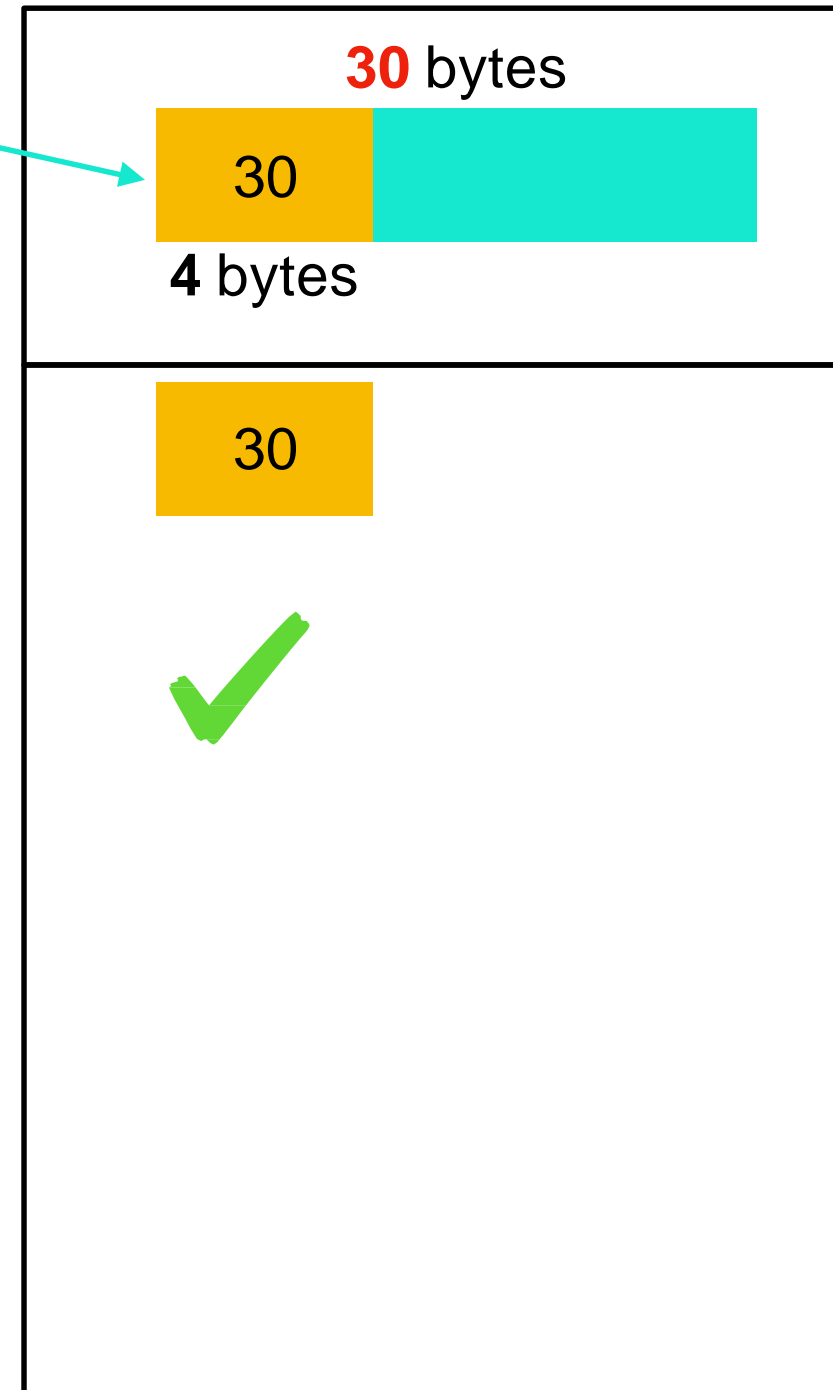


Why Double-Fetch?

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9     return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size < PERF_ATTR_SIZE_VER0)
14    return -EINVAL;
```

```
24 memcpy(out, attr, attr->size);
```

Adapted from `perf_copy_attr` in file `kernel/events/core.c`




Why Double-Fetch?

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9     return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size < PERF_ATTR_SIZE_VER0)
14    return -EINVAL;
15
16  // second fetch
17  if (copy_from_user(attr, uattr, size))
18    return -EFAULT;
```

```
24 memcpy(buf, attr, attr->size);
```

Adapted from `perf_copy_attr` in file `kernel/events/core.c`




Why Double-Fetch?

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9     return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size < PERF_ATTR_SIZE_VER0)
14    return -EINVAL;
15
16  // second fetch
17  if (copy_from_user(attr, uattr, size))
18    return -EFAULT;
19
20  . . . . .
21 }
```

```
24 memcpy(buf, attr, attr->size);
```


Adapted from `perf_copy_attr` in file `kernel/events/core.c`



What Goes Wrong in This Process?

Up-until First-Fetch

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9     return -EFAULT;
```



```
24 memcpy(out, attr, attr->size);
```

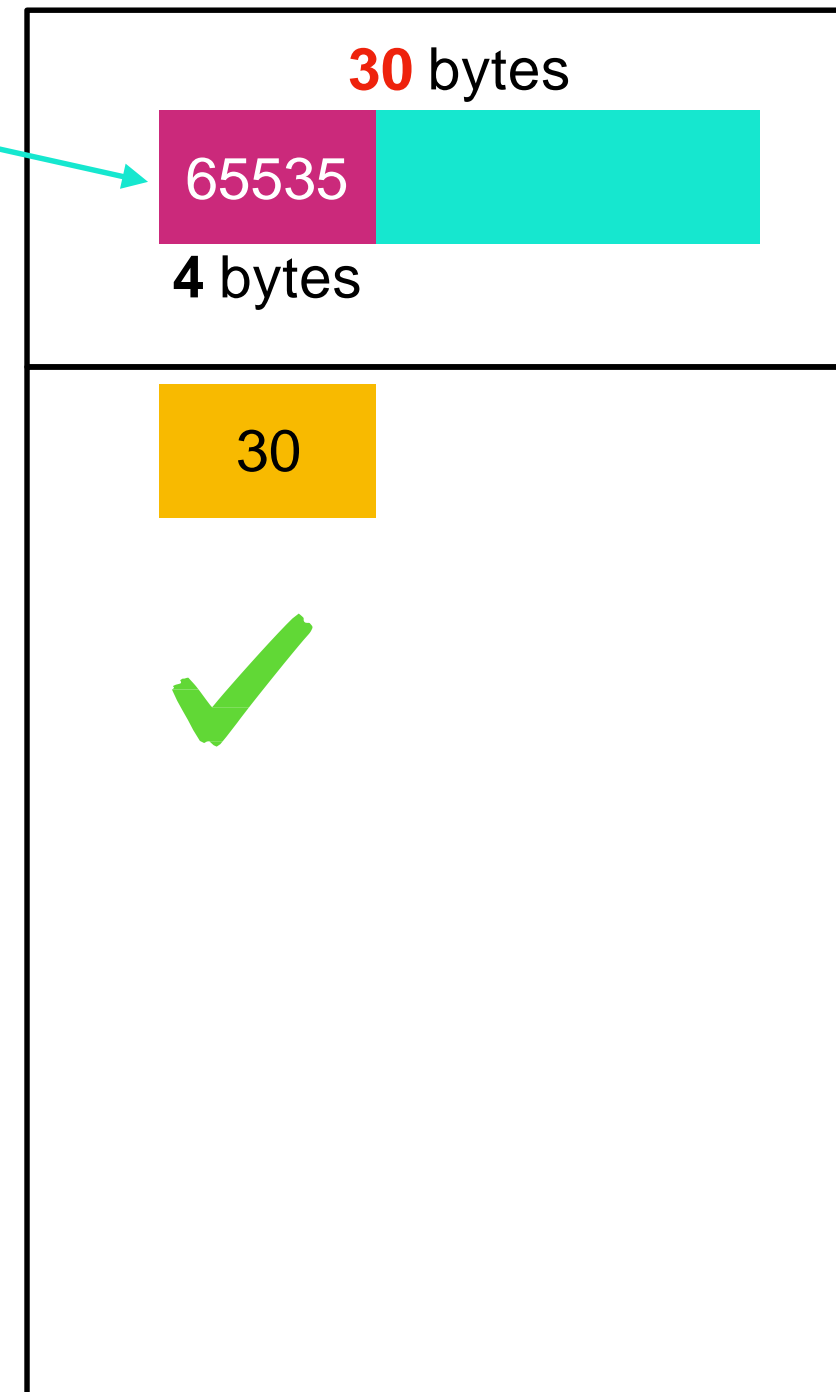
Adapted from `perf_copy_attr` in file `kernel/events/core.c`

Wrong Assumption: Atomicity in Syscall

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9     return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size <
14      PERF_ATTR_SIZE_VER0)
```

```
24 memcpy(buf, attr, attr->size);
```

Adapted from `perf_copy_attr` in file `kernel/events/core.c`

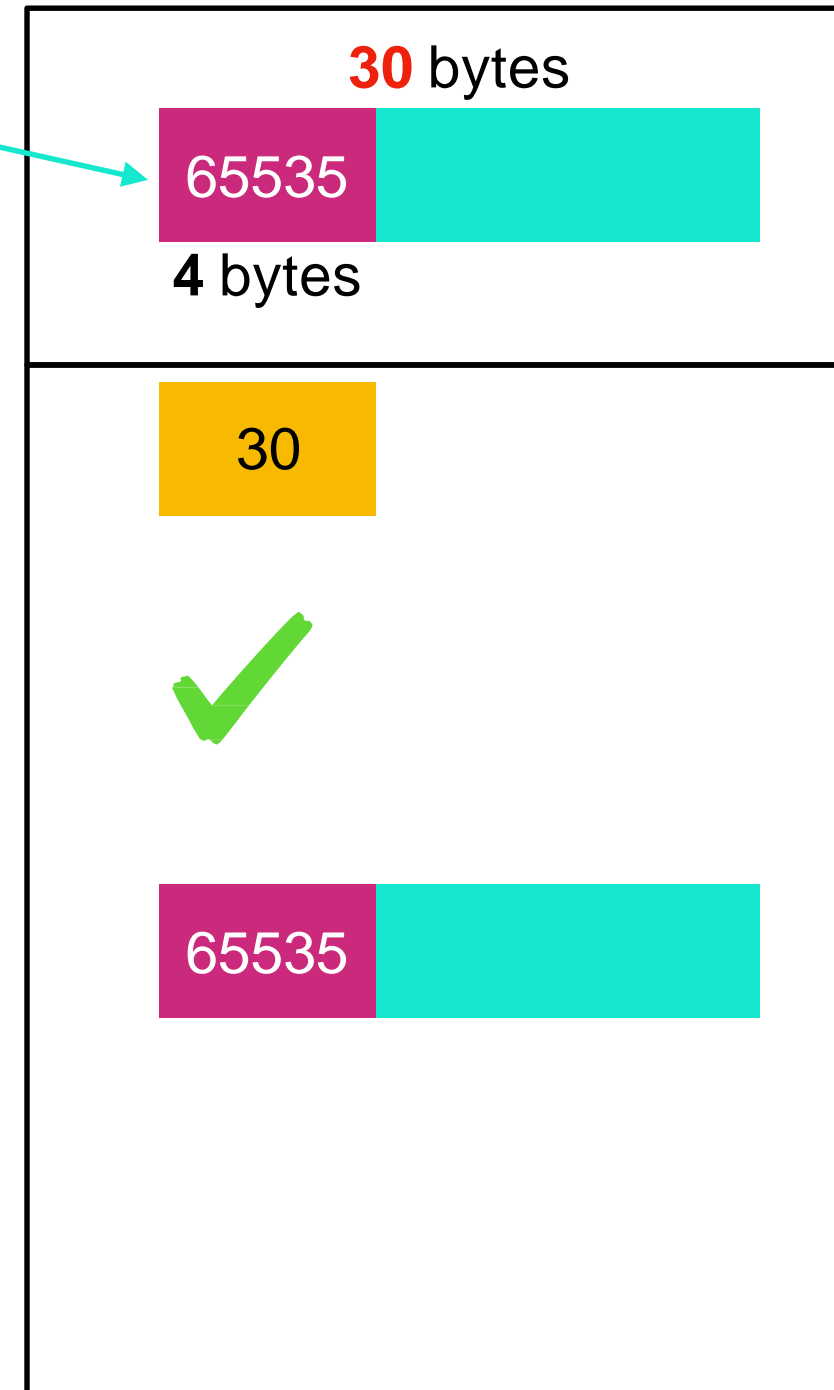


Wrong Assumption: Atomicity in Syscall

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9     return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size < PERF_ATTR_SIZE_VER0)
14    return -EINVAL;
15
16  // second fetch
17  if (copy_from_user(attr, uattr, size))
18    return -EFAULT;
19
20  . . . . .
21 }
```

```
24 memcpy(buf, attr, attr->size);
```

Adapted from `perf_copy_attr` in file `kernel/events/core.c`



When The Exploit Happens


```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9       return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size < PERF_ATTR_SIZE_VER0)
14      return -EINVAL;
15
16  // second fetch
17  if (copy_from_user(attr, uattr, size))
18      return -EFAULT;
19
20  . . . . .
21 }
22
23 // BUG: when attr->size is used
    later
24 copy_to_user(ubuf, attr, attr-
    >size);
```

Adapted from `perf_copy_attr` in file `kernel/events/core.c`

30 bytes



30



kernel information leak!

Double-Fetch Bugs: Towards A Formal Definition

Fetch: A pair (A, S) , where

A - the starting address of the fetch,

S - the size of memory copied into kernel.

Overlapped-fetch: Two fetches, (A_0, S_0) and (A_1, S_1) , where


$$A_0 \leq A_1 < A_0 + S_0 \quad // \quad A_1 \leq A_0 < A_1 + S_1$$

- The overlapped memory region is marked as (A_{01}, S_{01}) .
- The copied value during 1st fetch is $(A_{01}, S_{01}, 0)$
- The copied value during 2nd fetch is $(A_{01}, S_{01}, 1)$.

Overlapped-Fetch Case 1

`get_user(attr, &uptr->attr)`

`copy_from_user(kptr, uptr, size)`




$(A_{01}, S_{01}, 0)$ attr

$(A_{01}, S_{01}, 1)$ kptr->attr

Overlapped-Fetch Case 2

```
copy_from_user(  
    khdr, uptr, sizeof(struct hdr)  
)
```

```
copy_from_user(  
    kmsg, uptr, khdr->size  
)
```



$(A_{01}, S_{01}, 0)$ khdr->size, khdr->type, ...

$(A_{01}, S_{01}, 1)$ kmsg->size, kmsg->type, ...

Double-Fetch Bugs: Towards A Formal Definition

Control dependence: A variable $V \in (A_{01}, S_{01})$ and V must satisfy a set of constraints before the second fetch can happen.

Double-Fetch Bugs: Towards A Formal Definition

Control dependence: A variable $V \in (A_{01}, S_{01})$ and V must satisfy a set of constraints before the second fetch can happen.

```
1 void tls_setsockopt_simplified(char __user *arg) {
2     struct tls_crypto_info header, *full = /* allocated before */;
3
4     // first fetch
5     if (copy_from_user(&header, arg, sizeof(struct tls_crypto_info)))
6         return -EFAULT;
7
8     // protocol check
9     if (header.version != TLS_1_2_VERSION)
10        return -ENOTSUPP;
11
12    // second fetch
13    if (copy_from_user(full, arg,
14        sizeof(struct tls12_crypto_info_aes_gcm_128)))
15        return -EFAULT;
16
17    // BUG: full->version might not be TLS_1_2_VERSION
18    do_sth_with(full);
19 }
```

The diagram illustrates the control flow of the code. A red arrow points from the first fetch (line 5) to the protocol check (line 9). A blue arrow points from the protocol check (line 9) to the second fetch (line 13). A red arrow also points from the first fetch (line 5) to the second fetch (line 13), indicating a direct path that bypasses the protocol check.

Overlapped variable V :
header.version

The constraint it must satisfy:
header.version == TLS_1_2_VERSION

Expect:
full->version == TLS_1_2_VERSION

Double-Fetch Bugs: Towards A Formal Definition

Data dependence: A variable $V \in (A_{01}, S_{01})$ and V is consumed before or on the second fetch (e.g., involved in calculation, passed to function calls, etc).

Double-Fetch Bugs: Towards A Formal Definition

Data dependence: A variable $V \in (A_{01}, S_{01})$ and V is consumed before or on the second fetch.

```
1 void mptctl_simplified(unsigned long arg) {
2     mpt_ioctl_header khdr, __user *uhdr = (void __user *) arg;
3     MPT_ADAPTER *iocp = NULL;
4
5     // first fetch
6     if (copy_from_user(&khdr, uhdr, sizeof(khdr)))
7         return -EFAULT;
8
9     // dependency lookup
10    if (mpt_verify_adapter(khdr.iocnum, &iocp) <
11        return -EFAULT;
12
13    // dependency usage
14    mutex_lock(&iocp->iocctl_cmds.mutex);
15    struct mpt_fw_xfer kfwdl, __user *ufwdl = (void __user *) arg;
16
17    // second fetch
18    if (copy_from_user(&kfwdl, ufwdl, sizeof(struct mpt_fw_xfer)))
19        return -EFAULT;
20
21    // BUG: kfwdl.iocnum might not equal to khdr.iocnum
22    mptctl_do_fw_download(kfwdl.iocnum, .....);
23    mutex_unlock(&iocp->iocctl_cmds.mutex);
24 }
```

Overlapped variable V :

`khdr.iocnum`

Data dependence:

`mpt_verify_adapter(khdr.iocnum, &iocp)`

Expect:

`kfwdl.iocnum == khdr.iocnum`

Double-Fetch Bugs: Towards A Formal Definition

1. Two fetches from userspace memory that cover an **overlapped** region.
2. A relation must exist on the overlapped region between the two fetches. The relation can be either **control-dependence** or **data-dependence**.
3. We cannot **prove** that the relation established after first fetch still holds after the second fetch.

If all conditions are satisfied: a user thread might race condition to change the content in the overlapped region, and thus, to destroy the relation.

How to Find Double-Fetch Bugs?

How to Find Double-Fetch Bugs?

1. Find as many double-fetch pairs as possible, construct the code paths associated with each pair.
2. Symbolically check each code path and determine whether the two fetches makes a double-fetch bug.

Fetch Pair Collection

Goal: Statically enumerate all pairs of fetches that could possibly occur.

Fetch Pairs Collection

Goal: Statically enumerate all pairs of fetches that could possibly occur.

Ideal solution (top-down):

- ✓ 1. Identify all fetches in the kernel
- ✗ 2. Construct a complete, inter-procedural CFG for the whole kernel
- ✗ 3. Perform pair-wise reachability tests for each pair of fetches

Fetch Pairs Collection

Goal: Statically enumerate all pairs of fetches that could possibly occur.

Ideal solution (top-down):

- ✓ 1. Identify all fetches in the kernel
- ✗ 2. Construct a complete, inter-procedural CFG for the whole kernel
- ✗ 3. Perform pair-wise reachability tests for each pair of fetches

Our solution (bottom-up):

- ✓ 1. Identify all fetches in the kernel
- ✓ 2. For each fetch, within the function it resides in, scan its reaching instructions for fetches or fetch-involved functions

Bottom-up Fetch Pairs Collection

```
static void enclosing_function(  
    struct msg_hdr__user *uptr,  
    struct msg_full *kptr  
) {  
    ...  
    ...  
    ...  
    ...  
    ...  
    ...  
    ...  
    ...  
    ...  
    if (copy_from_user(kptr, uptr, size))  
        return -EFAULT;  
    ...  
}
```

Start from a fetch →

Bottom-up Fetch Pairs Collection

```
static void enclosing_function(  
    struct msg_hdr__user *uptr,  
    struct msg_full *kptr  
) {
```

...
...
...
...
...
...
...

Search through the
reaching instructions



```
    if (copy_from_user(kptr, uptr, size))  
        return -EFAULT;  
    ...  
}
```

Bottom-up Fetch Pairs Collection

[Case 1]
Found another fetch
==>
found a fetch pair

```
static void enclosing_function(  
    struct msg_hdr__user *uptr,  
    struct msg_full *kptr  
) {  
    ...  
    ...  
    if (get_user(size, &uptr->size))  
        return -EFAULT;  
    ...  
    ...  
    if (copy_from_user(kptr, uptr, size))  
        return -EFAULT;  
    ...  
}
```

Bottom-up Fetch Pairs Collection

[Case 2]
Found a fetch-involved
function
==>
inline the function,
found a fetch pair

```
static void enclosing_function(  
    struct msg_hdr__user *uptr,  
    struct msg_full *kptr  
) {  
    ...  
    ...  
    size = get_size_from_user(uptr);  
    ...  
    ...  
    ...  
    ...  
    if (copy_from_user(kptr, uptr, size))  
        return -EFAULT;  
    ...  
}
```

Bottom-up Fetch Pairs Collection

[Case 3]
No fetch-related
instruction
==>
Not a double-fetch



```
static void enclosing_function(  
    struct msg_hdr__user *uptr,  
    struct msg_full *kptr  
) {  
    ...  
    ...  
    ...  
    ...  
    ...  
    ...  
    ...  
    if (copy_from_user(kptr, uptr, size))  
        return -EFAULT;  
    ...  
}
```

How to Find Double-Fetch Bugs?

- ✓ 1. Find as many double-fetch pairs as possible, construct the code paths associated with each pair.
- 2. Symbolically check each code path and determine whether the two fetches makes a double-fetch bug.

Symbolic Checking

Goal: Symbolically execute the code path that connects two fetches and determine whether the two fetches satisfy all the criteria set in formal definition of double-fetch bug, i.e.

- Overlapp
- Have a relation (control or data dependence)
- We cannot prove the relation still holds after second fetch

Symbolic Checking

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9     return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size < PERF_ATTR_SIZE_VER0)
14    return -EINVAL;
15
16  // second fetch
17  if (copy_from_user(attr, uattr, size))
18    return -EFAULT;
19
20  . . . . .
21 }
22
23 // BUG: when attr->size is used later
24 memcpy(buf, attr, attr->size);
```

UI)

Symbolic Checking

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr __user *uattr,
3    struct perf_event_attr *attr) {
4
5     u32 size;
6
7     // first fetch
8     if (get_user(size, &uattr->size))
9         return -EFAULT;
10
11    // sanity checks
12    if (size > PAGE_SIZE ||
13        size < PERF_ATTR_SIZE_VER0)
14        return -EINVAL;
15
16    // second fetch
17    if (copy_from_user(attr, uattr, size))
18        return -EFAULT;
19
20    .....
21 }
22
23 // BUG: when attr->size is used later
24 memcpy(buf, attr, attr->size);
```

```
1 // init root SR
2 $0 = PARM(0), @0 = UMEM(0) // uattr
3 $1 = PARM(1), @1 = KMEM(1) // attr
4 - - -
```

UI)

Symbolic Checking

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5     u32 size;
6
7     // first fetch
8     if (get_user(size, &uattr-
9 >size))
10        return -EFAULT;
11    // sanity checks
12    if (size > PAGE_SIZE ||
13        size < PERF_ATTR_SIZE_VER0)
14        return -EINVAL;
15
16    // second fetch
17    if (copy_from_user(attr, uattr, size))
18        return -EFAULT;
19
20    . . . . .
21 }
22
23 // BUG: when attr->size is used later
24 memcpy(buf, attr, attr->size);
```

```
1 // init root SR
2 $0 = PARM(0), @0 = UMEM(0) // uattr
3 $1 = PARM(1), @1 = KMEM(1) // attr
4 - - -
5 // first fetch
6 fetch(F1): {A = $0 + 4, S = 4}
7 $2 = @0(4, 7, U0), @2 = nil // size
8 - - -
```

U1)

Symbolic Checking

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9       return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size <
14      PERF_ATTR_SIZE_VER0)
15      return -EINVAL;
16  // second fetch
17  if (copy_from_user(attr, uattr,
18                      sizeof(*attr))
19      return -EFAULT;
20
21  . . . . .
22
23  // BUG: when attr->size is used later
24  memcpy(buf, attr, attr->size);
```

```
1 // init root SR
2 $0 = PARM(0), @0 = UMEM(0) // uattr
3 $1 = PARM(1), @1 = KMEM(1) // attr
4 - - -
5 // first fetch
6 fetch(F1): {A = $0 + 4, S = 4}
7 $2 = @0(4, 7, U0), @2 = nil // size
8 - - -
9 // sanity checks
10 assert $2 <= PAGE_SIZE
11 assert $2 >= PERF_ATTR_SIZE_VER0
12 - - -
```

U1)

Symbolic Checking

```

1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9       return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size < PERF_ATTR_SIZE_VER0)
14      return -EINVAL;
15
16  // second fetch
17  if (copy_from_user(attr, uattr,
18                      size)
19      return -EFAULT;
20
21  . . . . .
22
23  // BUG: when attr->size is used later
24  memcpy(buf, attr, attr->size);

```

```

1 // init root SR
2 $0 = PARM(0), @0 = UMEM(0) // uattr
3 $1 = PARM(1), @1 = KMEM(1) // attr
4 - - -
5 // first fetch
6 fetch(F1): {A = $0 + 4, S = 4}
7 $2 = @0(4, 7, U0), @2 = nil // size
8 - - -
9 // sanity checks
10 assert $2 <= PAGE_SIZE
11 assert $2 >= PERF_ATTR_SIZE_VER0
12 - - -
13 // second fetch
14 fetch(F2): {A = $0, S = $2}
15 @1(0, $2 - 1, K) = @0(0, $2 - 1, U1)
16 - - -

```

U1)

Symbolic Checking

```

1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9       return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size < PERF_ATTR_SIZE_VER0)
14      return -EINVAL;
15
16  // second fetch
17  if (copy_from_user(attr, uattr, size))
18      return -EFAULT;
19
20  . . . . .
21 }
22
23 // BUG: when attr->size is used later
24 memcpy(buf, attr, attr->size);

```

```

1 // init root SR
2 $0 = PARM(0), @0 = UMEM(0) // uattr
3 $1 = PARM(1), @1 = KMEM(1) // attr
4 ---
5 // first fetch
6 fetch(F1): {A = $0 + 4, S = 4}
7 $2 = @0(4, 7, U0), @2 = nil // size
8 ---
9 // sanity checks
10 assert $2 <= PAGE_SIZE
11 assert $2 >= PERF_ATTR_SIZE_VER0
12 ---
13 // second fetch
14 fetch(F2): {A = $0, S = $2}
15 @1(0, $2 - 1, K) = @0(0, $2 - 1, U1)
16 ---
17 // check fetch overlap
18 assert F2.A <= F1.A < F2.A + F2.S
19      OR F1.A <= F2.A < F1.A + F1.S
20 [solve]
21 --> satisfiable with @0(4, 7, U)

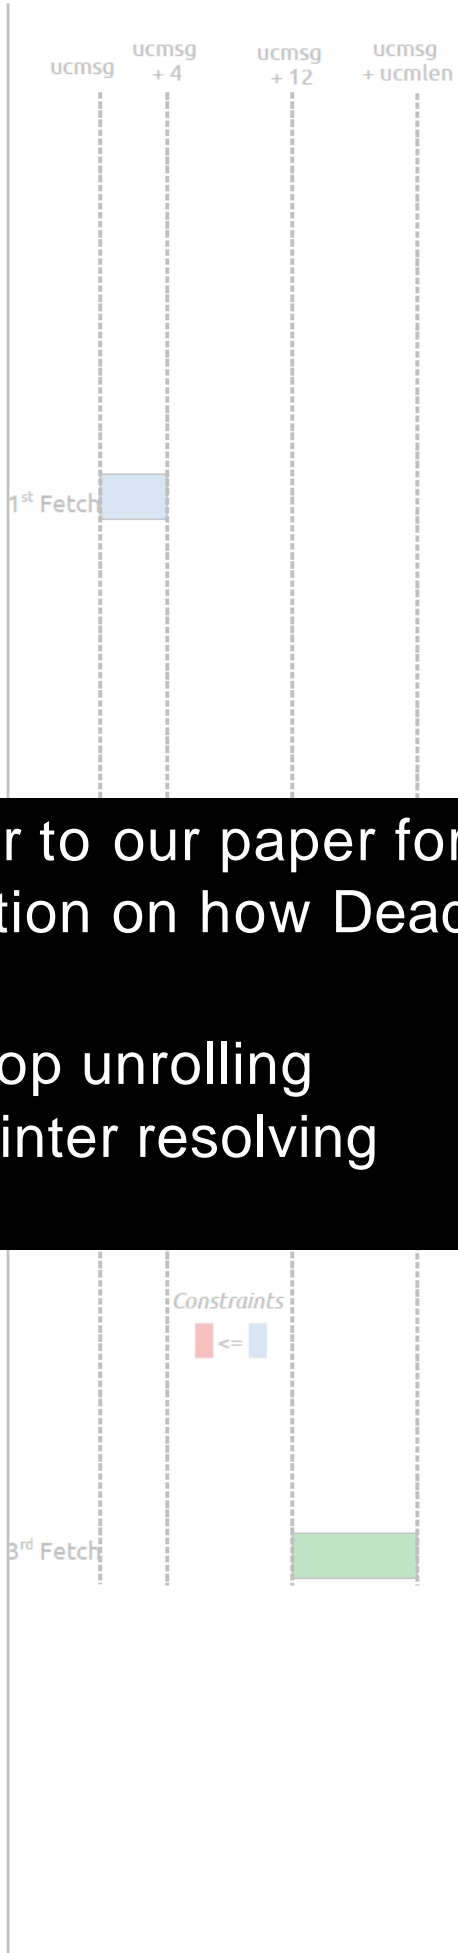
```

U1)

Symbolic Checking

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr _user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
6
7   // first fetch
8   if (get_user(size, &uattr->size))
9       return -EFAULT;
10
11  // sanity checks
12  if (size > PAGE_SIZE ||
13      size < PERF_ATTR_SIZE_VER0)
14      return -EINVAL;
15
16  // second fetch
17  if (copy_from_user(attr, uattr, size))
18      return -EFAULT;
19
20  . . . . .
21 }
22
23 // BUG: when attr->size is used later
24 memcpy(buf, attr, attr->size);
```

```
1 // init root SR
2 $0 = PARM(0), @0 = UMEM(0) // uattr
3 $1 = PARM(1), @1 = KMEM(1) // attr
4 ---
5 // first fetch
6 fetch(F1): {A = $0 + 4, S = 4}
7 $2 = @0(4, 7, U0), @2 = nil // size
8 ---
9 // sanity checks
10 assert $2 <= PAGE_SIZE
11 assert $2 >= PERF_ATTR_SIZE_VER0
12 ---
13 // second fetch
14 fetch(F2): {A = $0, S = $2}
15 @1(0, $2 - 1, K) = @0(0, $2 - 1, U1)
16 ---
17 // check fetch overlap
18 assert F2.A <= F1.A < F2.A + F2.S
19      OR F1.A <= F2.A < F1.A + F1.S
20 [solve]
21   --> satisfiable with @0(4, 7, U)
22 // check double-fetch bug
23 [prove] @0(4, 7, U0) == @0(4, 7, U1)
24   --> fail: no constraints on @0(4, 7, U1)
```

<pre> 1 int cmsghdr_from_user_compat_to_kern 2 (struct msghdr *kmsg, char *kbuf) { 3 4 struct compat_cmsghdr __user *ucmsg; 5 compat_size_t ucmlen; 6 struct cmsghdr *kcmsg; 7 __kernel_size_t kcmllen, tmp; 8 9 // 1st loop: calculate message length 10 kcmllen = 0; 11 ucmsg = kmsg->msg_control; 12 while (ucmsg != NULL) { 13 // first batch of fetches 14 if (get_user(ucmlen, &ucmsg->cmsg_len)) 15 return -EFAULT; 16 17 tmp = ucmlen + sizeof(struct cmsghdr) 18 - sizeof(struct compat_cmsghdr); 19 20 kcmllen += tmp; 21 ucmsg = (char *)ucmsg + ucmlen; 22 } 23 24 // 2nd loop: copy the whole message 25 kcmsg = kbuf; 26 ucmsg = kmsg->msg_control; 27 while (ucmsg != NULL) { 28 // second batch of fetches 29 if (get_user(ucmlen, &ucmsg->cmsg_len)) 30 return -EFAULT; 31 32 tmp = ucmlen + sizeof(struct cmsghdr) 33 - sizeof(struct compat_cmsghdr); 34 35 // sanity check, but insufficient 36 if (kbuf + kcmllen - (char *)kcmsg < tmp) 37 return -EINVAL; 38 39 // irrelevant fetch 40 if (copy_from_user(41 (char *)kcmsg + sizeof(*kcmsg), 42 (char *)ucmsg + sizeof(*ucmsg), 43 (ucmlen - sizeof(*ucmsg)))) 44 return -EFAULT; 45 46 kcmsg = (char *)kcmsg + tmp; 47 ucmsg = (char *)ucmsg + ucmlen; 48 } 49 50 // BUG: the actual message length != kcmllen 51 kmsg->msg_controllen = kcmllen; 52 return 0; 53 } </pre>		<pre> 1 // init root SR 2 \$0 = \$PARM(0), @0 = \$KMEM(0) // kmsg 3 \$1 = \$PARM(1), @1 = \$KMEM(1) // kbuf 4 --- 5 // prepare for the 1st batch of fetches 6 \$2 = 0, @2 = nil // kcmllen_0 7 \$3 = @0(48, 55, K), @3 = \$UMEM(0) // ucmsg_0 8 --- 9 // unroll 1st loop 10 assert \$2 != NULL 11 fetch(F1) is {A = \$3 + 0, S = 4} 12 \$4 = @3(0, 3, U0), @4 = nil // ucmlen_0 13 \$5 = \$4 - 12 + 16, @5 = nil // tmp_0 14 \$6 = \$2 + \$5, @6 = nil // kcmllen_1 15 \$7 = \$3 + \$4, @7 = \$UMEM(1) // ucmsg_1 16 assert \$7 == NULL (i.e., @7 = nil) // exit loop 17 --- 18 // prepare for the 2nd batch of fetches 19 \$8 = \$1 @8 = \$KMEM(1) // kcmsg_0 20 \$9 = @0(48, 55, K) == \$3, @9 = @3 // ucmsg_2 21 --- 22 // unroll 2nd loop 23 24 } 25 @0 = nil // ucmlen_1 26 @1 = nil // tmp_1 27 28 > @3(0, 3, U0) >= @3(0, 3, U1) 29 30 \$10 - 12} 31 \$10 - 13, U0) 32 33 @2 = \$KMEM(2) // kcmsg_1 34 @3 = \$UMEM(3) // ucmsg_3 35 36 assert \$13 == NULL (i.e., @13 = nil) // exit loop 37 --- 38 // check fetch overlap 39 assert F2.A <= F1.A < F2.A + F2.S 40 AND F1.A <= F2.A < F1.A + F1.S 41 // --> satisfiable with @3(0, 3, U) 42 43 assert F3.A <= F1.A < F3.A + F3.S 44 AND F1.A <= F3.A < F1.A + F1.S 45 // --> unsatisfiable 46 47 assert F3.A <= F2.A < F3.A + F3.S 48 AND F2.A <= F3.A < F2.A + F2.S 49 // --> unsatisfiable 50 51 // check double-fetch bug 52 prove @3(0, 3, U0) == @3(0, 3, U1) 53 // --> fail, as @3(0, 3, U0) >= @3(0, 3, U1) </pre>
(a) C source code	(b) Memory access patterns	(c) Symbolic representation and checking

Please refer to our paper for a comprehensive demonstration on how Deadline handles

1. Loop unrolling
2. Pointer resolving

Findings

- 24 bugs found in total
 - 23 bugs in Linux kernel and 1 in FreeBSD kernel
- 9 bugs have been patched with the fix we provide
- 4 bugs are acknowledged, we are still working on the fix
- 9 bugs are pending for review
- 2 bugs are marked as “won’t fix”

Double-Fetch Bug Mitigations

- The basic idea is to re-assure the control-dependence and data-dependence between the two fetches. In other words, the **automaticity** in user space memory fetches during the execution of the syscall.

Double-Fetch Bug Mitigations

- The basic idea is to re-assure the control-dependence and data-dependence between the two fetches. In other words, the **automaticity** in user space memory fetches during the execution of the syscall.
- Based on our experience and our communications with kernel developers, we found **four patterns** in patching double-fetch bugs.

Double-Fetch Bug Mitigations

1. Override after second fetch.

```
1 kernel/events/core.c | 2 ++
2 1 file changed, 2 insertions(+)
3
4 diff --git a/kernel/events/core.c b/kernel/events/core.c
5 index ee20d4c..c0d7946 100644
6 --- a/kernel/events/core.c
7 +++ b/kernel/events/core.c
8 @@ -9611,6 +9611,8 @@ static int perf_copy_attr(struct perf_event_attr __user *uattr,
9     if (ret)
10         return -EFAULT;
11
12 + attr->size = size;
13 +
14     if (attr->__reserved_1)
15         return -EINVAL;
```

Override the overlapped memory (`attr->size`)
with the value from the first fetch (`size`).

Double-Fetch Bug Mitigations

2. Abort on change detected.

```
1 net/compat.c | 7 ++++++
2 1 file changed, 7 insertions(+)
3
4 diff --git a/net/compat.c b/net/compat.c
5 index 6ded6c8..2238171 100644
6 --- a/net/compat.c
7 +++ b/net/compat.c
8 @@ -185,6 +185,13 @@ int cmsghdr_from_user_compat_to_kern(struct msghdr *kmsg, struct sock *sk,
9         ucmsg = cmsg_compat_nxthdr(kmsg, ucmsg, ucmlen);
10     }
11
12 + /*
13 +  * check the length of messages copied in is the same as the
14 +  * what we get from the first loop
15 +  */
16 + if ((char *)kcmsg - (char *)kcmsg_base != kcmllen)
17 +     goto Eival;
18 +
19 /* Ok, looks like we made it. Hook it up and return success. */
20 kmsg->msg_control = kcmsg_base;
21 kmsg->msg_controllen = kcmllen;
```

Compare the new message length (kcmsg - kcmsg_base) with the value from the first fetch (kcmllen).

Double-Fetch Bug Mitigations

3. Refactor overlapped copies into incremental copies.

```
1  block/scsi_ioctl.c | 8 ++++++--
2  1 file changed, 7 insertions(+), 1 deletion(-)
3
4  diff --git a/block/scsi_ioctl.c b/block/scsi_ioctl.c
5  index 7440de4..8fe1e05 100644
6  --- a/block/scsi_ioctl.c
7  +++ b/block/scsi_ioctl.c
8  @@ -463,7 +463,13 @@ int sg_scsi_ioctl(struct request_queue *q, struct gendisk *disk, fmode_t mode,
9      */
10     err = -EFAULT;
11     req->cmd_len = cmdlen;
12 -   if (copy_from_user(req->cmd, sic->data, cmdlen))
13 +
14 +   /*
15 +    * avoid copying the opcode twice
16 +    */
17 +   memcpy(req->cmd, &opcode, sizeof(opcode));
18 +   if (copy_from_user(req->cmd + sizeof(opcode),
19 +                       sic->data + sizeof(opcode), cmdlen - sizeof(opcode)))
20       goto error;
21
22   if (in_len && copy_from_user(buffer, sic->data + cmdlen, in_len))
```

When copying the whole message, skip the information copied in the first fetch (+ sizeof(opcode)).

Double-Fetch Bug Mitigations

4. Refactor overlapped copies into a single-fetch.


```

1 drivers/isdn/i4l/isdn_ppp.c | 37 ++++++-----
2 1 file changed, 25 insertions(+), 12 deletions(-)
3
4 diff --git a/drivers/isdn/i4l/isdn_ppp.c b/drivers/isdn/i4l/isdn_ppp.c
5 index 6c44609..cd2b3c6 100644
6 --- a/drivers/isdn/i4l/isdn_ppp.c
7 +++ b/drivers/isdn/i4l/isdn_ppp.c
8 @@ -825,7 +825,6 @@ isdn_ppp_write(int min, struct file *file, const char __user *buf, int count)
9     isdn_net_local *lp;
10     struct ipp_struct *is;
11     int proto;
12 -    unsigned char protobuf[4];
13
14     is = file->private_data;
15
16 @@ -839,24 +838,28 @@ isdn_ppp_write(int min, struct file *file, const char __user *buf, int count)
17     if (!lp)
18         printk(KERN_DEBUG "isdn_ppp_write: lp == NULL\n");
19     else {
20 -        /*
21 -         * Don't reset huptimer for
22 -         * LCP packets. (Echo requests).
23 -         */
24 -        if (copy_from_user(protobuf, buf, 4))
25 -            return -EFAULT;
26 -        proto = PPP_PROTOCOL(protobuf);
27 -        if (proto != PPP_LCP)
28 -            lp->huptimer = 0;
29 +        if (lp->isdn_device < 0 || lp->isdn_channel < 0) {
30 +            unsigned char protobuf[4];
31 +            /*
32 +             * Don't reset huptimer for
33 +             * LCP packets. (Echo requests).
34 +             */
35 +            if (copy_from_user(protobuf, buf, 4))
36 +                return -EFAULT;
37 +
38 +            proto = PPP_PROTOCOL(protobuf);
39 +            if (proto != PPP_LCP)
40 +                lp->huptimer = 0;
41
42 -        if (lp->isdn_device < 0 || lp->isdn_channel < 0)
43 -            return 0;
44 +    }
45
46     if ((dev->drv[lp->isdn_device]->flags & DRV_FLAG_RUNNING) &&
47         lp->dialstate == 0 &&
48         (lp->flags & ISDN_NET_CONNECTED)) {
49         unsigned short hl;
50         struct sk_buff *skb;
51 +        unsigned char *cpy_buf;
52         /*
53          * we need to reserve enough space in front of
54          * sk_buff. old call to dev_alloc_skb only reserved
55 @@ -869,11 +872,21 @@ isdn_ppp_write(int min, struct file *file, const char __user *buf, int count)
56         return count;
57     }
58     skb_reserve(skb, hl);
59 -    if (copy_from_user(skb_put(skb, count), buf, count))
60 +    cpy_buf = skb_put(skb, count);
61 +    if (copy_from_user(cpy_buf, buf, count))
62     {
63         kfree_skb(skb);
64         return -EFAULT;
65     }
66 +
67 +    /*
68 +     * Don't reset huptimer for
69 +     * LCP packets. (Echo requests).
70 +     */
71 +    proto = PPP_PROTOCOL(cpy_buf);
72 +    if (proto != PPP_LCP)
73 +        lp->huptimer = 0;
74 +
75     if (is->debug & 0x40) {
76         printk(KERN_DEBUG "ppp xmit: len %d\n", (int) skb->len);
77         isdn_ppp_frame_log("xmit", skb->data, skb->len, 32, is->unit, lp->ppp_slot);

```

Such a strategy is usually very complex and requires careful refactoring.

Limitations of Deadline

- **Source code coverage**
 - Files not compilable under LLVM.
 - Special combination of kernel configs (e.g., `CONFIG_*`).
- **Execution path construction**
 - Limit on total number of paths explored per fetch pair (4096).
 - Loop unrolling (limited to unroll once only).
- **Symbolic checking**
 - Ignores inline assemblies.
 - Imprecise pointer to memory object mapping.
 - Assumption on enclosing function.

Conclusion

- Detecting double-fetch bugs without a precise and formal definition has led to many false alerts and tremendous manual effort.
- Deadline is based on a precise modeling of double-fetch bugs and achieves both high accuracy and high scalability.
- Application beyond kernels: hypervisors, browsers, TEE, etc.
- Logic bugs are on the rise! We hope that more logic bugs can be modeled and checked systematically.

<https://github.com/sslabs-gatech/deadline>