

kAFL : Hardware-Assisted Feedback Fuzzing for OS Kernels

This paper is included in the Proceedings of the 26th USENIX Security
Symposium

What is fuzzing?

1. insert an input file into a program
2. mutate the input (randomly?)
3. let the program open the input file
4. execute to see where is crashed

What is Feedback Fuzzing?


uses mechanisms to learn which inputs are interesting and which are not

The difficulty of fuzzing kernel

1. kernel-level code has significantly more non-determinism than the average ring 3 program—mostly due to interrupts, kernel threads, statefulness, and similar mechanisms.
2. crashes and timeouts mandate the use of virtualization to be able to catch faults and continue gracefully
3. close source

Basic Knowledge

- Intel CPU Privilege: ring 0, ring 1, ring 2, ring3. (ring 0 and ring 3 in windows). ring 0 -> kernel, ring 3 -> user space
- Intel VT-x
 - physical CPU
 - logical CPU
 - vCPU
 - VM
 - VMM



Basic Knowledge

- Intel PT

With the fifth generation of Intel Coreprocessors, Intel PT provides execution and branch tracing information.


general execution information and control flow information packets.

- Taken-Not-Taken(TNT): e.g. jmp
- Target IP (TIP): e.g. ret
- Flow Update Packets(FUP): interrupts and traps

- CR3 filter

distinguish kernel mode or user mode

Overview of KAFL



1. fuzzing logic
2. VM infrastructure
3. user mode agent (loader and agent)

Procedure

1. loader uses HC_SUBMIT_PANIC function to transfer panic handler address to QEMU-PT, in order to get information before crashing (modify the panic handler).
2. loader uses HC_GET_PROGRAM to request actual user agent, in order to initialize the fuzzing.
3. user agent uses HC_SUBMIT_CR3 function to KVM-PT to let VM give current CR3 to QEMU-PT; Finally trigger HC_SUBMIT_BUFFER to tell host the input location, then finish the initialization.
4. during the main loop, user agent uses HC_GET_INPUT function to get inputs.
5. fuzzer generates inputs to QEMU-PT, then QEMU-PT puts the inputs to the specific address, then trigger VM-ENTRY VM's execution;
6. VM-ENTRY also triggers PT tracer.
7. during fuzzing, QEMU-PT decodes tracing data, once get back to the user space, VM send signal HC_FINISHED, then trigger VM-EXIT to exit the VM.
8. The result is passed to the logic for further processing

KVM-PT

Why KVM-PT?

- monitor vCPU instead of logical CPU

How to avoid unwanted trace data?

- MSR autoload capabilities of Intel VT-x

How to save tracing?

- ToPA

How to solve ToPA overflowing?

- add a small buffer of variable length

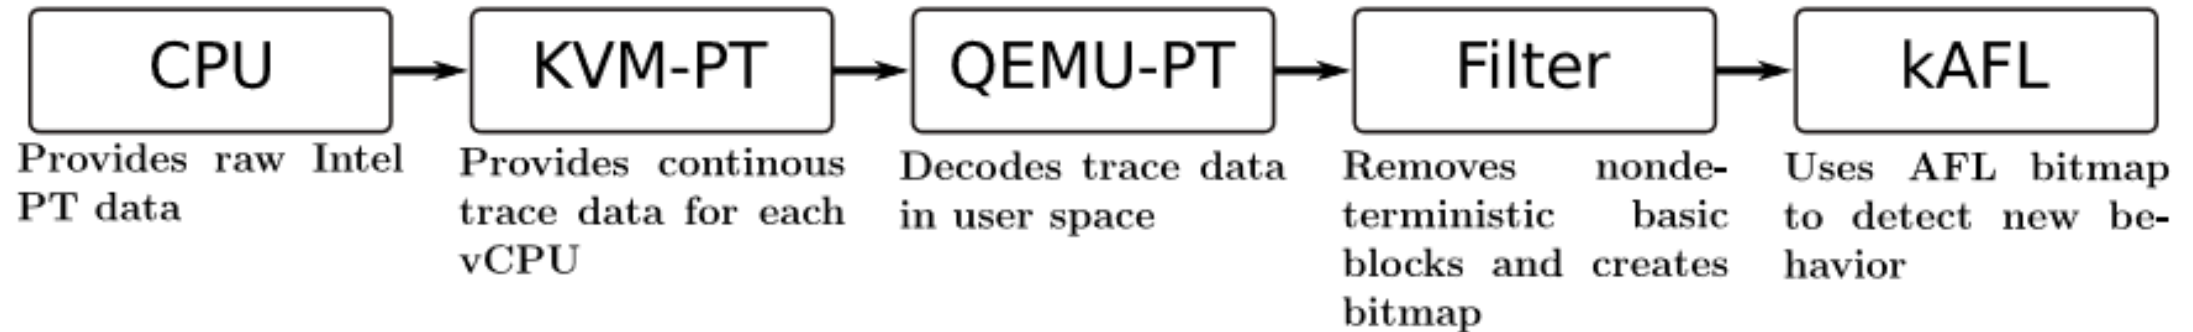
QEMU-PT

Why QEMU-PT?

- a user space counterpart of KVM-PT, enable, disable and configure Intel PT at run time

- decode the tracing data, and convert it to a bitmap

Procedure



Evaluation

Windows: NTFS Div-by-Zero

macOS: HFS Div-by-Zero

- HFS Assertion Fail

- HFS Use-After-Free

- APFS Memory Corruption

Linux: keyctl Null Pointer Dereference

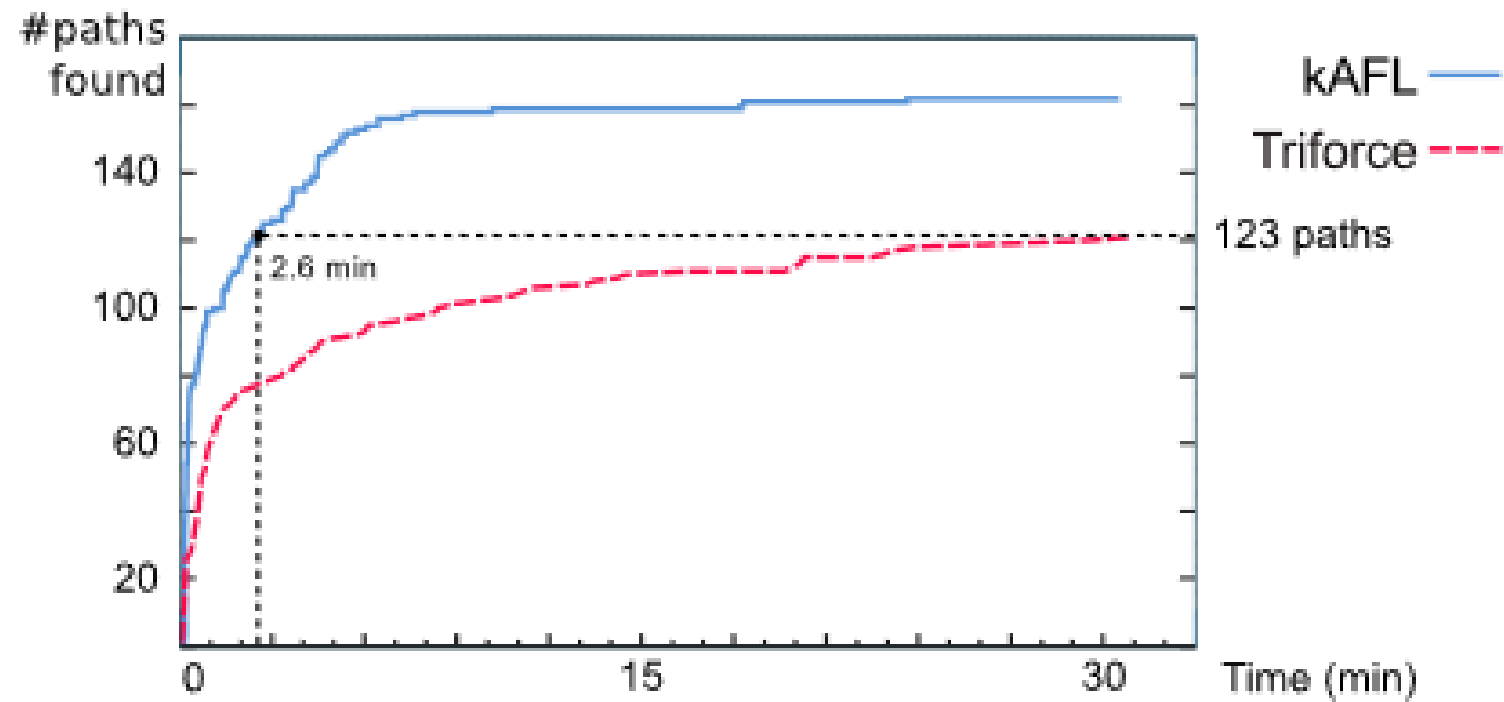
- ext4 Memory Corruption

- ext4 Error Handling


Evaluation

```
1  jsmn_parser parser;
2  jsmntok_t tokens[5];
3  jsmn_init(&parser);
4
5  int res=jsmn_parse(&parser, input, size, tokens, 5);
6  if(res >= 2){
7      if(tokens[0].type == JSMN_STRING){
8          int json_len = tokens[0].end - tokens[0].
              start;
9          int s = tokens[0].start;
10         if(json_len > 0 && input[s+0] == 'K'){
11             if(json_len > 1 && input[s+1] == 'A'){
12                 if(json_len > 2 && input[s+2] == 'F'){
13                     if(json_len > 3 && input[s+3] == 'L'){
14                         panic(KERN_INFO "KAFL...\n");
15                     }
16                 }
17             }
18         }
19     }
```

Evaluation



Evaluation



Limitation

- technique limitation
- decoding just-in-time code

APPENDDDIX

The Introduction of AFL

Overview

AFL(American Fuzzing Loop) is an open-source fuzzing tool.

- It can fuzz software in user space
- It can fuzz open-source software or closed-source software

The use of AFL

open-sourced software:

gcc -----> afl-gcc

g++-----> afl-g++

clang -----> afl-clang

e.g CC=/path/to/afl/afl-gcc ./configure

The use of AFL

start fuzzing:

stdin:

```
./afl-fuzz -i testcase_dir -o findings_dir /path/to/program [...params...]
```

file:

```
./afl-fuzz -i testcase_dir -o findings_dir /path/to/program @@
```

Example

```
root@hwsrv-521890:~/upx/src# afl-fuzz -i ../../afl_in -o ../../afl_out ./upx.out @@
```

american fuzzy lop 2.52b (upx.out)			
process timing		overall results	
run time : 0 days, 0 hrs, 0 min, 9 sec		cycles done : 0	
last new path : 0 days, 0 hrs, 0 min, 3 sec		total paths : 39	
last uniq crash : 0 days, 0 hrs, 0 min, 6 sec		uniq crashes : 4	
last uniq hang : none seen yet		uniq hangs : 0	
cycle progress		map coverage	
now processing : 0 (0.00%)		map density : 2.28% / 3.26%	
paths timed out : 0 (0.00%)		count coverage : 1.21 bits/tuple	
stage progress		findings in depth	
now trying : bitflip 1/1		favored paths : 1 (2.56%)	
stage execs : 5124/182k (2.81%)		new edges on : 32 (82.05%)	
total execs : 6849		total crashes : 7 (4 unique)	
exec speed : 735.1/sec		total tmouts : 0 (0 unique)	
fuzzing strategy yields		path geometry	
bit flips : 0/0, 0/0, 0/0		levels : 2	
byte flips : 0/0, 0/0, 0/0		pending : 39	
arithmetics : 0/0, 0/0, 0/0		pend fav : 1	
known ints : 0/0, 0/0, 0/0		own finds : 38	
dictionary : 0/0, 0/0, 0/0		imported : n/a	
havoc : 0/0, 0/0		stability : 100.00%	
trim : 0.00%/1412, n/a			
[cpu:191%]			

close-sourced: (poor efficiency)

compile quem_mode

start fuzzing:

```
afl_fuzz -i afl_in -o afl_out -Q program -a @@
```

Example

```
root@hwsrv-521890:~/afl-master# ./afl-fuzz -i ../afl_in_2 -o ../afl_out_2 -Q readelf -a @@
```

american fuzzy lop 2.52b (readelf)

process timing		overall results
run time : 0 days, 0 hrs, 42 min, 15 sec		cycles done : 0
last new path : 0 days, 0 hrs, 0 min, 1 sec		total paths : 380
last uniq crash : none seen yet		uniq crashes : 0
last uniq hang : none seen yet		uniq hangs : 0
cycle progress	map coverage	
now processing : 0 (0.00%)	map density : 2.49% / 4.93%	
paths timed out : 0 (0.00%)	count coverage : 2.10 bits/tuple	
stage progress	findings in depth	
now trying : bitflip 1/1	favored paths : 1 (0.26%)	
stage execs : 123k/135k (91.16%)	new edges on : 162 (42.63%)	
total execs : 127k	total crashes : 0 (0 unique)	
exec speed : 56.93/sec (slow!)	total tmouts : 0 (0 unique)	
fuzzing strategy yields	path geometry	
bit flips : 0/0, 0/0, 0/0	levels : 2	
byte flips : 0/0, 0/0, 0/0	pending : 380	
arithmetics : 0/0, 0/0, 0/0	pend fav : 1	
known ints : 0/0, 0/0, 0/0	own finds : 379	
dictionary : 0/0, 0/0, 0/0	imported : n/a	
havoc : 0/0, 0/0	stability : 100.00%	
trim : 0.00%/1042, n/a		
[cpu: 180%]		

Algorithm

1. Load user-supplied initial test cases into the queue,
2. Take next input file from the queue,
3. Attempt to trim the test case to the smallest size that doesn't alter the measured behavior of the program,
4. Repeatedly mutate the file using a balanced and well-researched variety of traditional fuzzing strategies,
5. If any of the generated mutations resulted in a new state transition recorded by the instrumentation, add mutated output as a new entry in the queue.
6. Go to 2.

Details

Why using afl-gcc/g++ instead of gcc/g++?

It will hook the branch instruction.

(after compiling the source code to assembly language, using afl-as)

Why hooking the branch instruction?

improving the coverage

calculating the count of hitting the branch

About Coverage

- Function
- Basic Block
- Edge

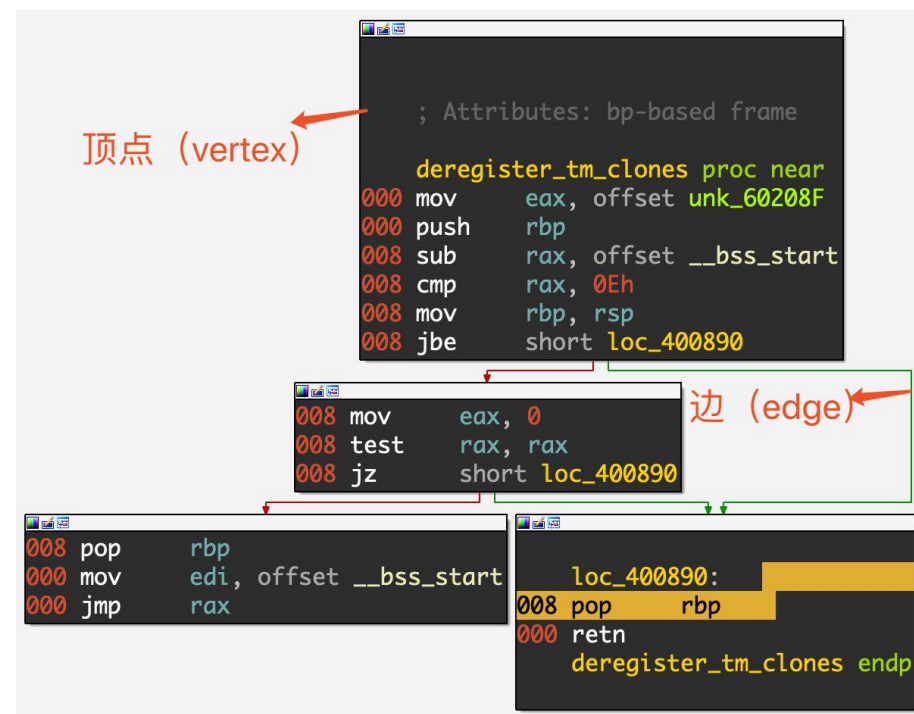
```
W = 0;  
x = x + y;  
if( x > z ){  
    y = x;  
    x++;  
} else {  
    y = z;  
    z++;  
}  
W = x + z;
```

```
W = 0;  
x = x + y;  
if( x > z ){
```

```
    y = x;  
    x++;
```

```
    y = z;  
    z++;
```

```
W = x + z;
```



About Coverage

- tuple (last block + current block)

e.g. A -> B -> C -> D -> E (tuples: AB, BC, CD, DE)

- how to save and find a new transition?

count of branch execution ---> 8 bits

using global map to save

- what is interesting?

a new one

the one from one bucket to another bucket

```
static const u8
count_class_lookup8[256] = {
    [0]          = 0,
    [1]          = 1,
    [2]          = 2,
    [3]          = 4,
    [4 ... 7]    = 8,
    [8 ... 15]   = 16,
    [16 ... 31]  = 32,
    [32 ... 127] = 64,
    [128 ... 255] = 128
};
```

Mutation

- bitflip
- arithmetic
- interest
- dictionary
- havoc
- splice

Fork server

- fuzzer -----> fork server
- copy on write